

# Containers

**CS 106L, Fall '21**

**How do we handle data in the STL?**

# Today's agenda

- What makes a container?
- The Standard Template Library (STL)
- Containers (Stanford vs STL)
  - Sequence
  - Associative
  - How do they actually work?

# How do we store data in programs?

What do we need in order to do so?

# Python and Stanford Library

Classic examples of data structures:




```
lst = []  
dictionary = {}  
hash_table = set()  
string = "adc"
```



```
Vector<int> lst;  
Map<int, int> dictionary;  
Set<int> hash_table;  
std::string str = "adc";
```


# What makes a container?

- Let's design one! We'll be making a queue.




# What makes a container?

- Let's design one! We'll be making a queue.




# What makes a container?

- Let's design one! We'll be making a queue.



# What makes a container?


- Let's design one! We'll be making a queue.






# What makes a container?

- Let's design one! We'll be making a queue.



# What makes a container?

- Let's design one! We'll be making a queue.



# This seems like a pretty common problem, right?

- What if I have to make a queue for a movie ticket line?
  - for a car wash?
  - for a burst of API requests?
- We need a place to store common definitions for data structures.

# Standard Template Library



**The core of modern C++!**

# What's in the STL?

**Containers**

**Iterators**

**Functions**

**Algorithms**

# What's in the STL?

Containers

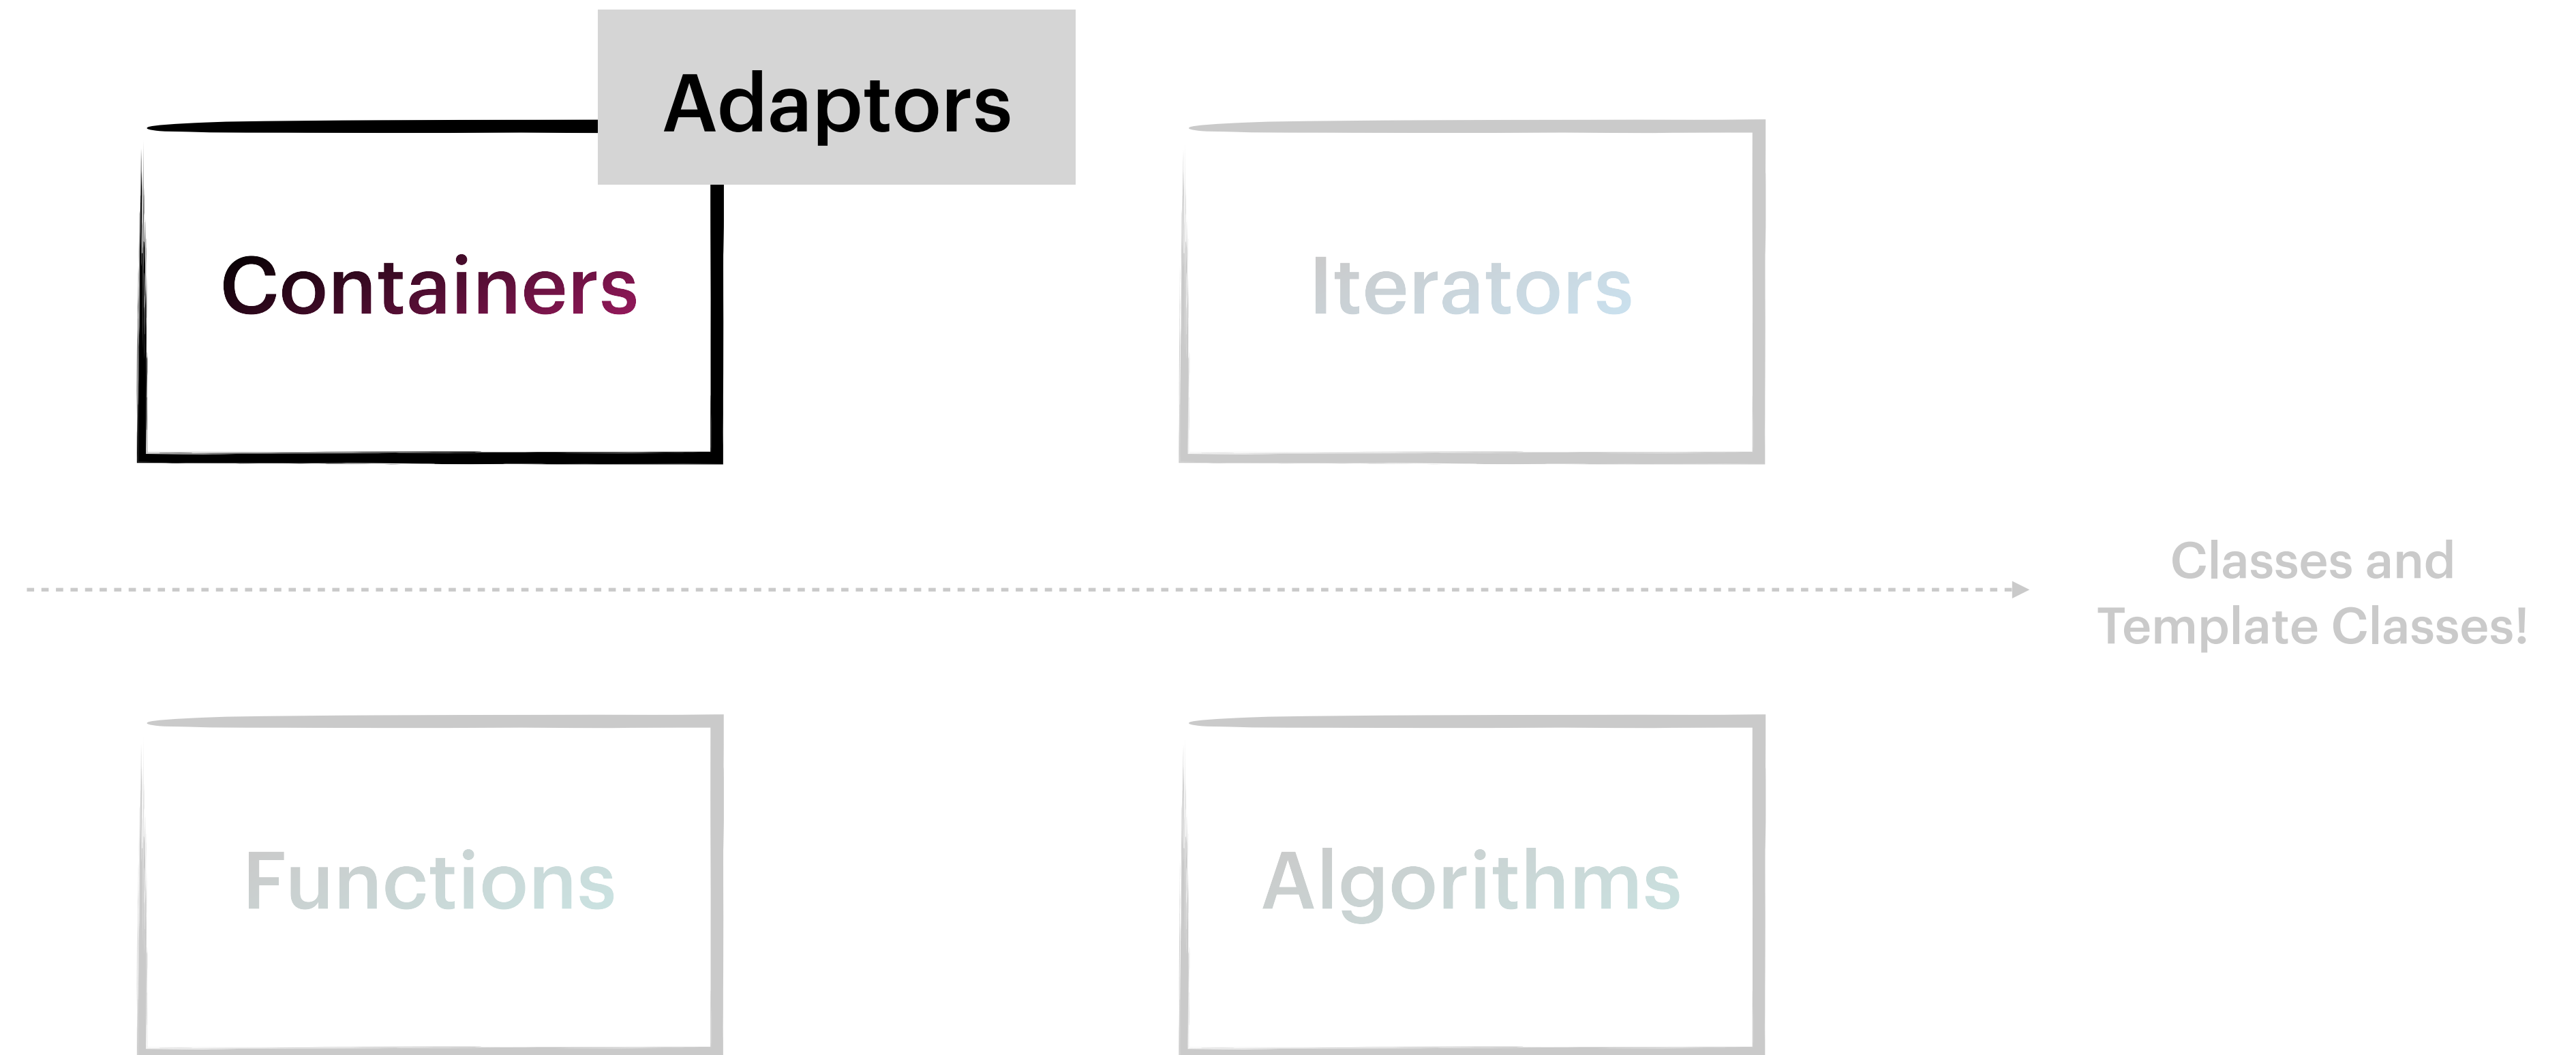
Iterators

Functions

Algorithms

Classes and  
Template Classes!

# What's in the STL?



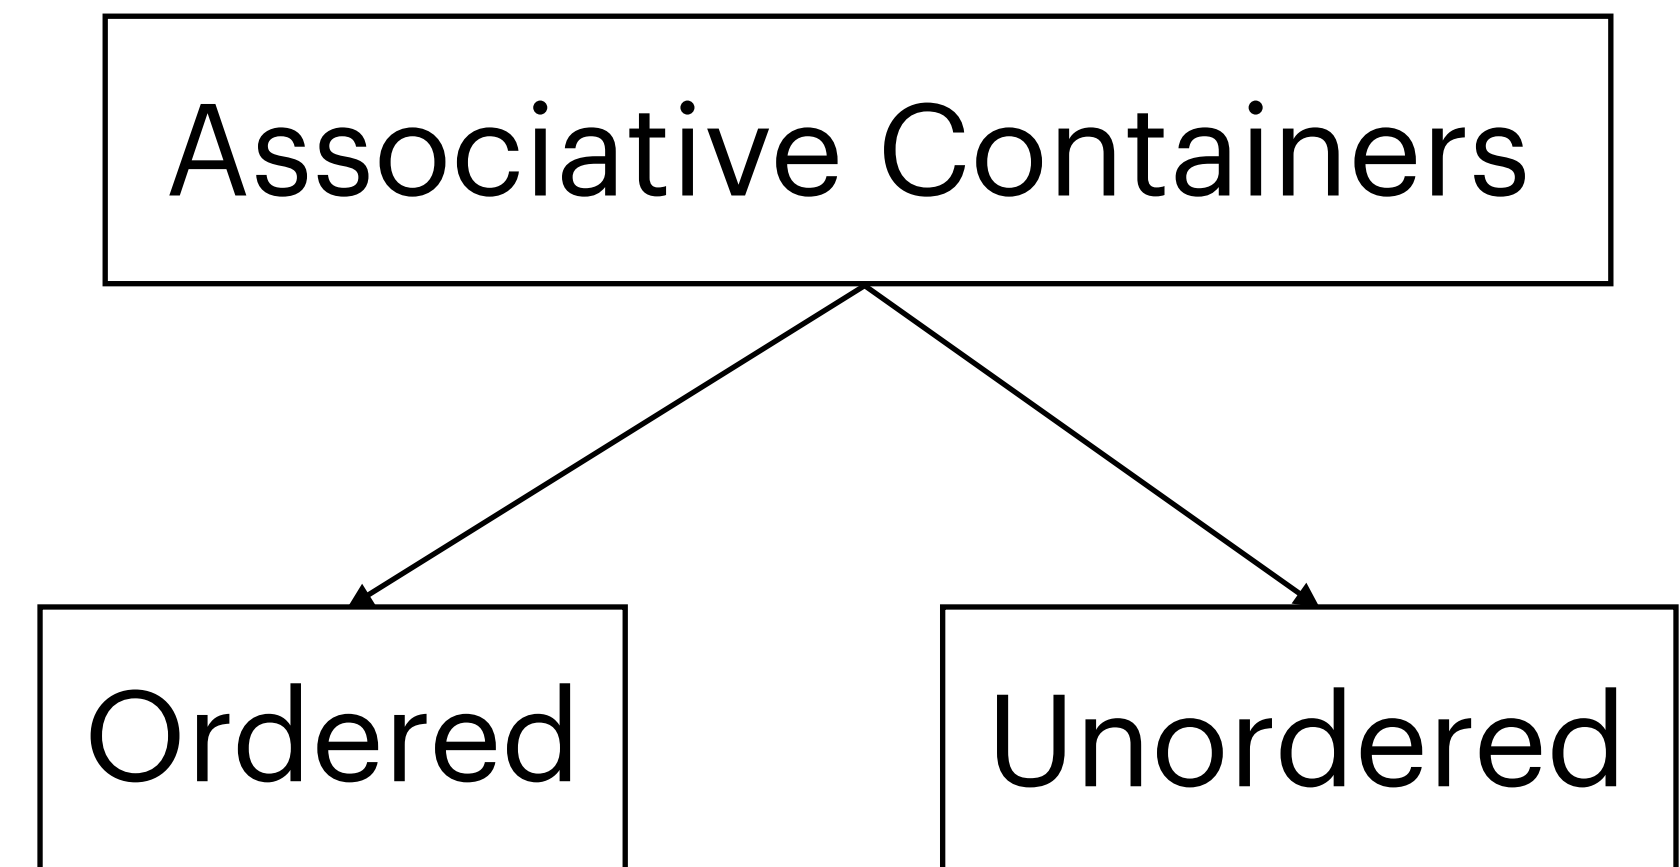

# Types of containers

- All containers can hold almost all elements.




# Types of containers

- All containers can hold almost all elements.



# Types of containers

- All containers can hold almost all elements.



<> vector (adding + removing elements at end)

↕ deque (adding + removing elements anywhere but end)

↓ list (adding + removing elements anywhere, but no random access)

() tuple (different data types, but immutable)

# How do we use the STL? (and an aside on “::”)

```
#import <vector>

int main () {
    std::vector<int> vec;
    ...
}
```

Just two steps!

1. Import the relevant STL feature
2. Use it with “std::<STL feature name here>”

“::” -> Scope Resolution Operator

For heavily used items, we can use certain classes and datatypes in the std namespace, e.g. “using std::string;”

```
using namespace std;
```

# Live Demo!

Let's compare the Stanford's Vector and STL's vector:  
(QT Creator Project)

# Stanford “Vector” vs STL “vector”

What you want to do	Stanford Vector<int>	std::vector<int>
Create a new, empty vector	Vector<int> vec;	std::vector<int> vec;
Create a vector with <b>n</b> copies of 0	Vector<int> vec(n);	std::vector<int> vec(n);
Create a vector with <b>n</b> copies of a value <b>k</b>	Vector<int> vec(n, k);	std::vector<int> vec(n, k);
Add a value <b>k</b> to the end of a vector	vec.add(k);	vec.push_back(k);
Remove all elements of a vector	vec.clear();	vec.clear();
Get the element at index <b>i</b>	int k = vec[i];	int k = vec[i]; (does not bounds check)
Check size of vector	vec.size();	vec.size();
Loop through vector by index <b>i</b>	for (int i = 0; i < vec.size(); ++i) ...	for (std::size_t i = 0; i < vec.size(); ++i) ...
Replace the element at index <b>i</b>	vec[i] = k;	vec[i] = k; (does not bounds check)


# Stanford “Vector” vs STL “vector”

What you want to do	Stanford <code>Vector&lt;int&gt;</code>	<code>std::vector&lt;int&gt;</code>
Insert k at some index i	<code>vec.insert(i, k);</code>	<code>vec.insert(vec.begin() + i, k);</code>
Remove the element at index i	<code>vec.remove(i);</code>	<code>vec.erase(vec.begin() + i);</code>
Get the sublist in range [i, j)	<code>v.subList(i, j);</code>	<code>std::vector&lt;int&gt; sum (vec.begin() + i, vec.begin() + j);</code>
Create a vector that is two vectors appended to each other	<code>Vector&lt;int&gt; v = v1 + v2;</code>	<code>// pretty complicated ngl</code>
Add j to the front of a vector	<code>vec.insert(0, i);</code>	<code>vec.insert(vec.begin(), k);</code>

We need ***iterators*** to understand these. Next Lecture!

# Types of containers

- In general, how do we pick between containers?



<> vector (adding + removing elements at end)

↕ deque (adding + removing elements anywhere but end)

↓ list (adding + removing elements anywhere, but no random access)

() tuple (different data types, but immutable)



# Live Demo!

Why can't we just use a vector for everything?  
`vector_time_trials.cpp`



# Why is there no `std::vector::push_front()`?


- Because it's super slow!
- This is what is happening to the elements inside the vector:



(in general, if something is extremely inefficient, C++ does not expose a method for it!)

# Types of containers

- In general, how do we pick between containers?



<> vector (adding + removing elements at end)

↕ deque (adding + removing elements anywhere but end)

↓ list (adding + removing elements anywhere, but no random access)

() tuple (different data types, but immutable)

# Live Demo!

Let's compare speeds of several operations!  
vector\_list\_deque.cpp

# How is a vector actually implemented?

- Internally, a **std::vector** consists of a fixed-sized array
  - It automatically resizes for you! (Arrays will be discussed in CS 106B after week 5)

# How is a vector actually implemented?

- Internally, a **std::vector** consists of a fixed-sized array
  - It automatically resizes for you! (Arrays will be discussed in CS 106B after week 5)

**size** = number of elements in a vector

**capacity** = amount of space saved for a vector

**size** = 6

**capacity** = 8

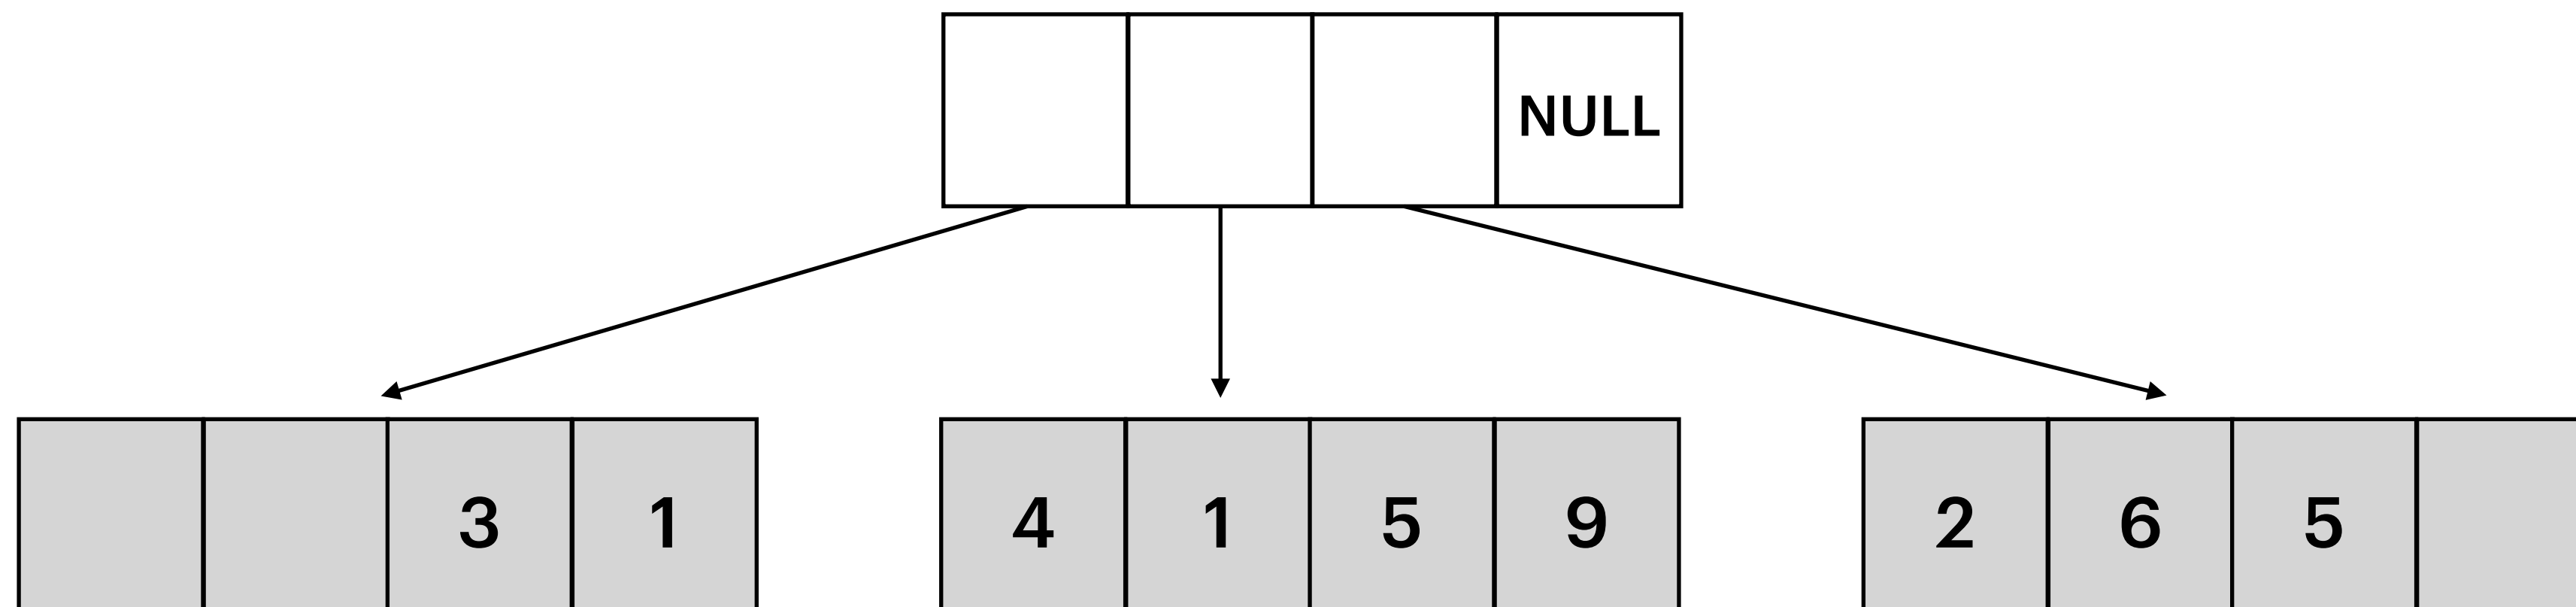
1	6	1	8	0	3		
---	---	---	---	---	---	--	--

# How is a deque actually implemented?

- There's no single, common implementation of a deque, but a common one looks like this:

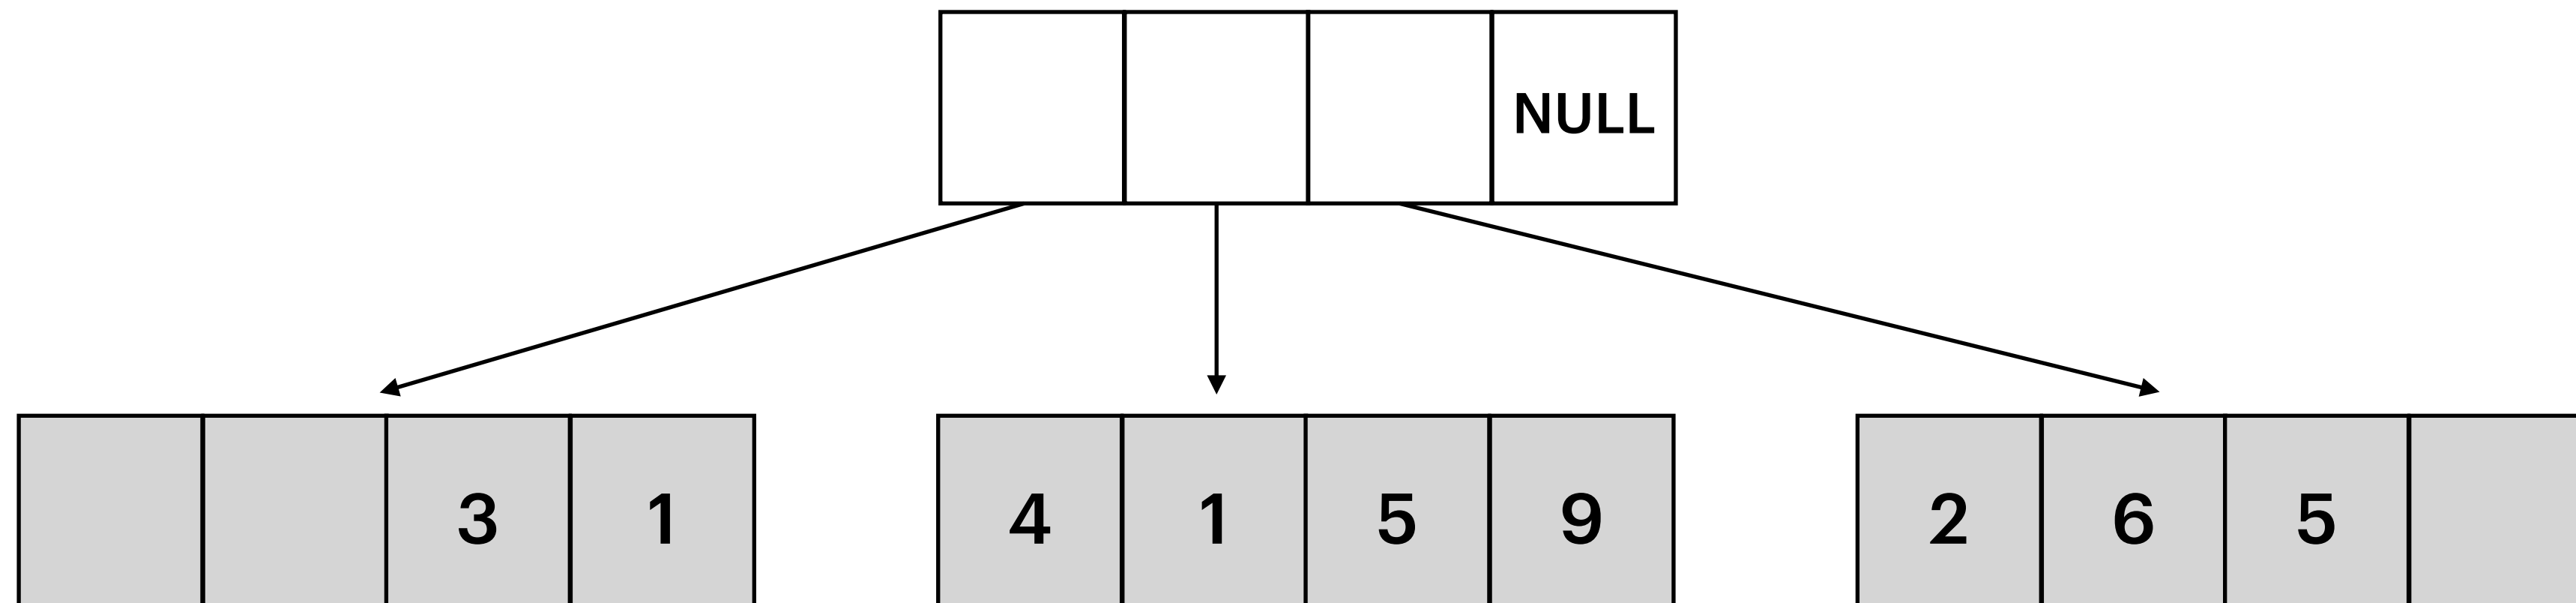
# How is a deque actually implemented?

- There's no single, common implementation of a deque, but a common one looks like this:



# How is a deque actually implemented?

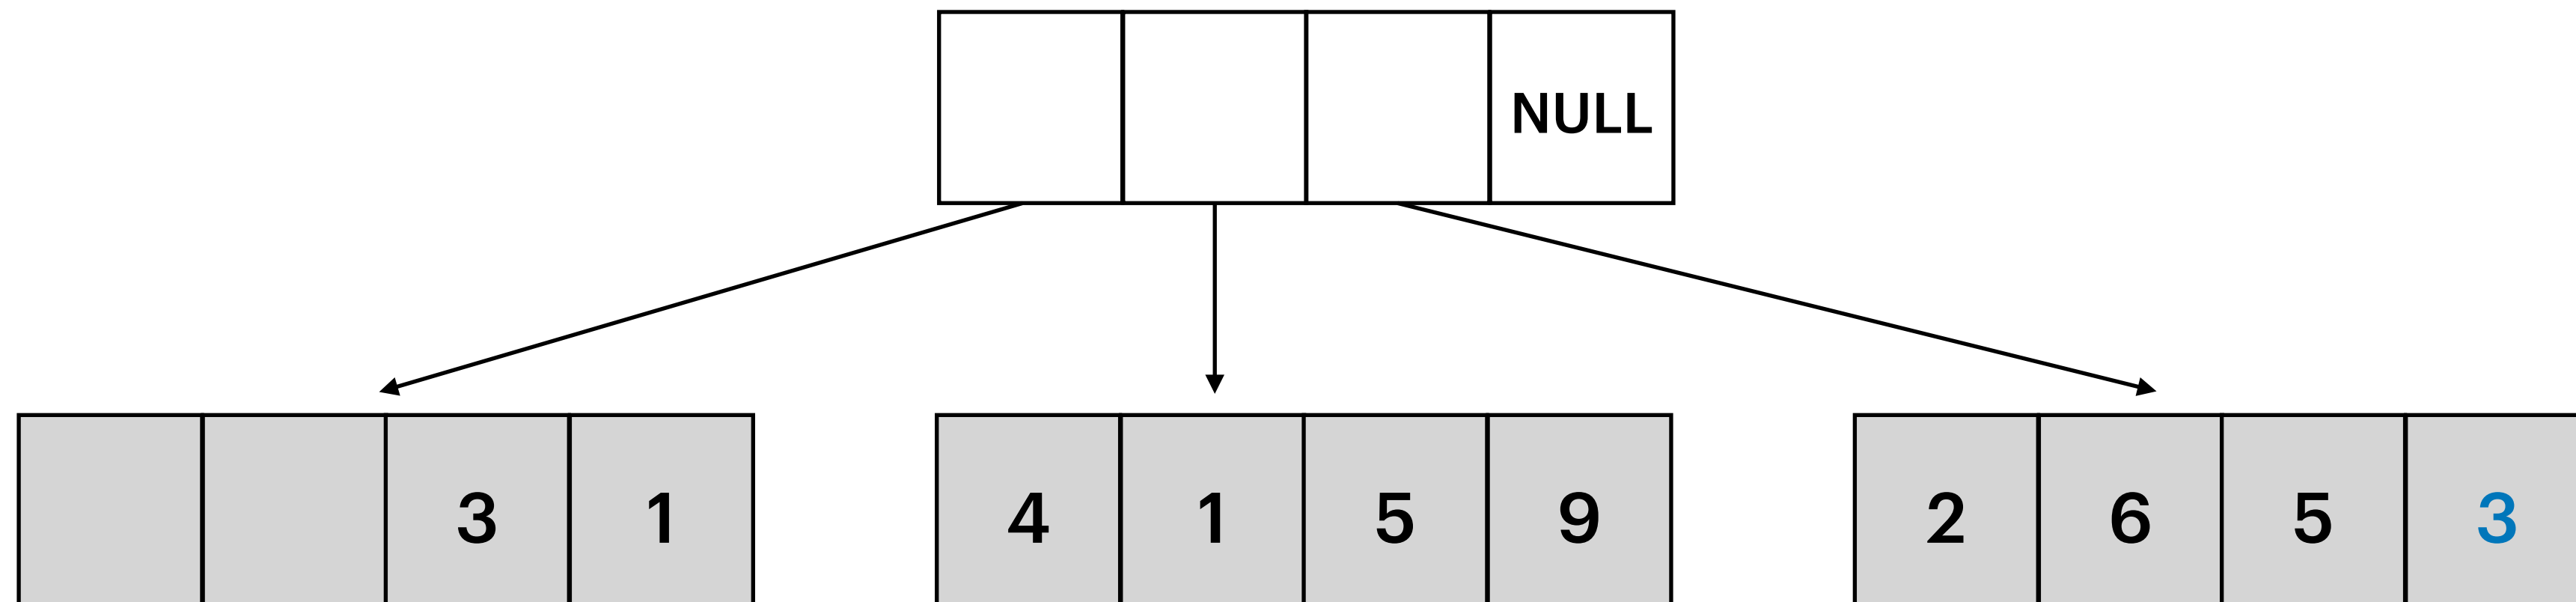
- So how do we **push\_back(3)**?





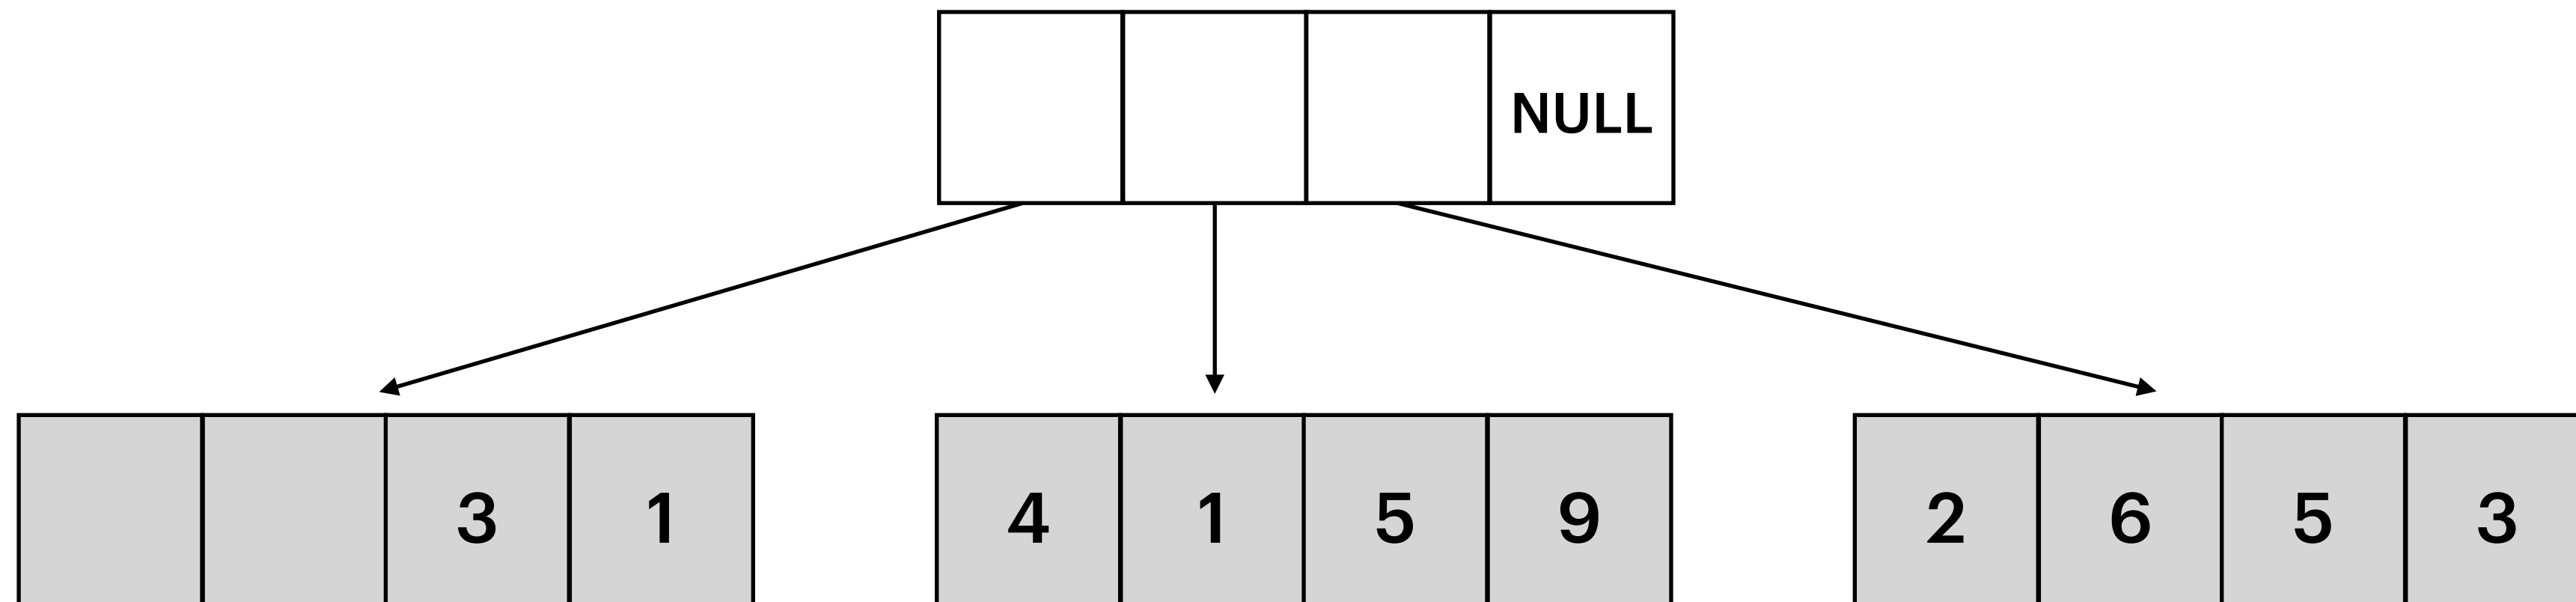
# How is a deque actually implemented?

- So how do we **push\_back(3)**?



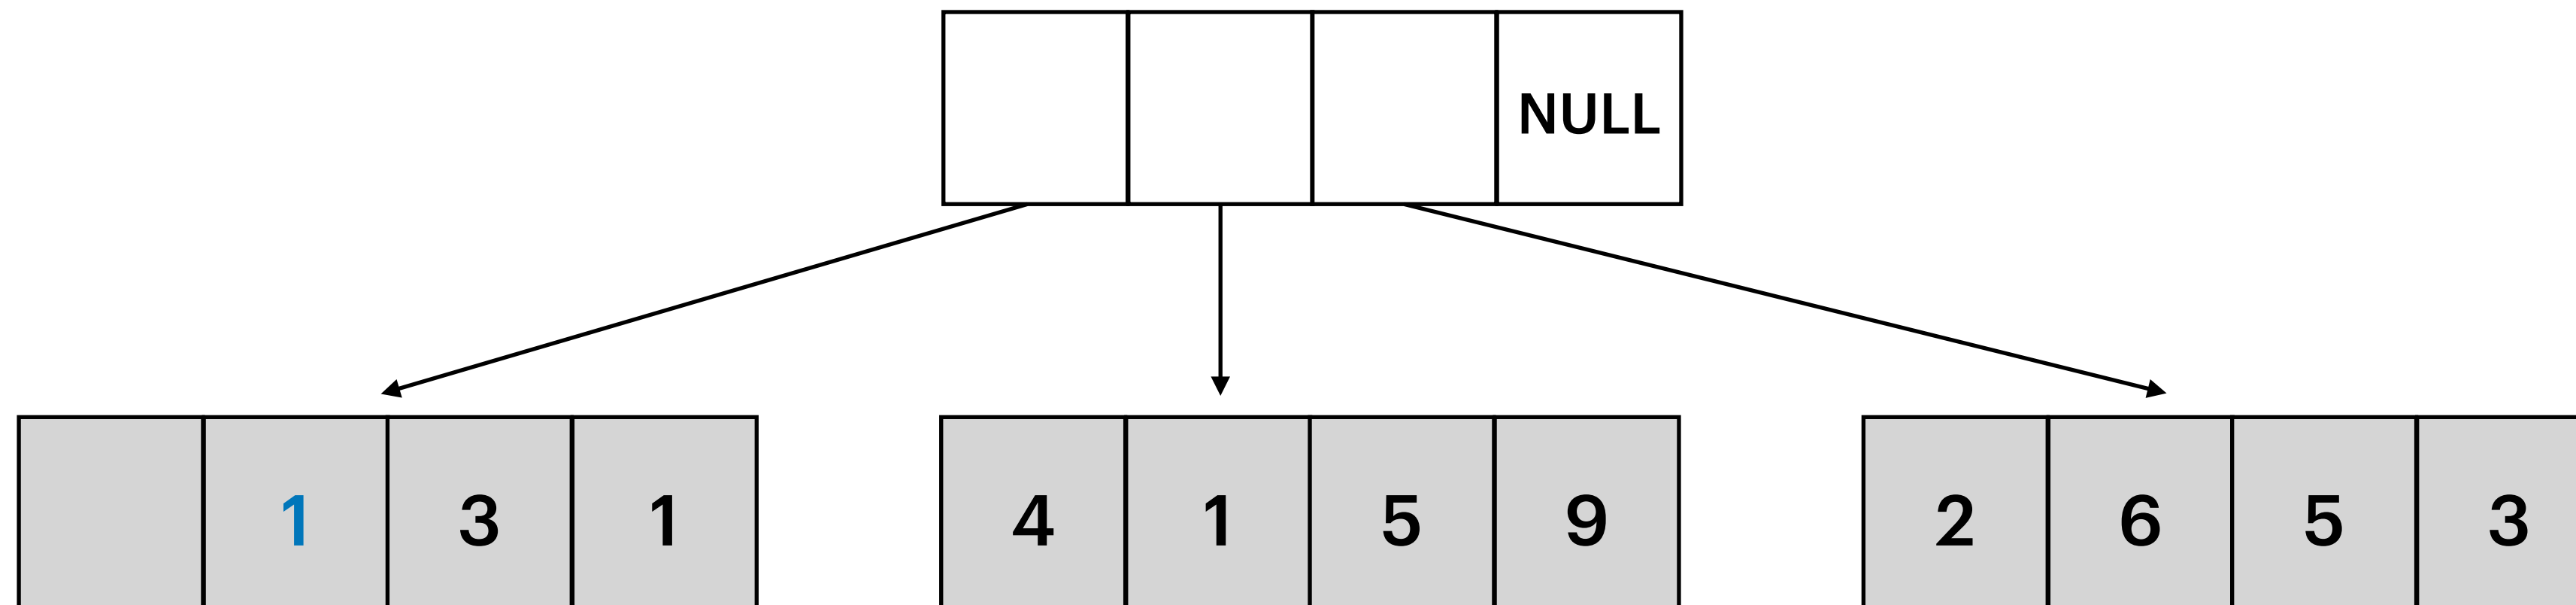
# How is a deque actually implemented?

- So how do we **push\_front(1)**?



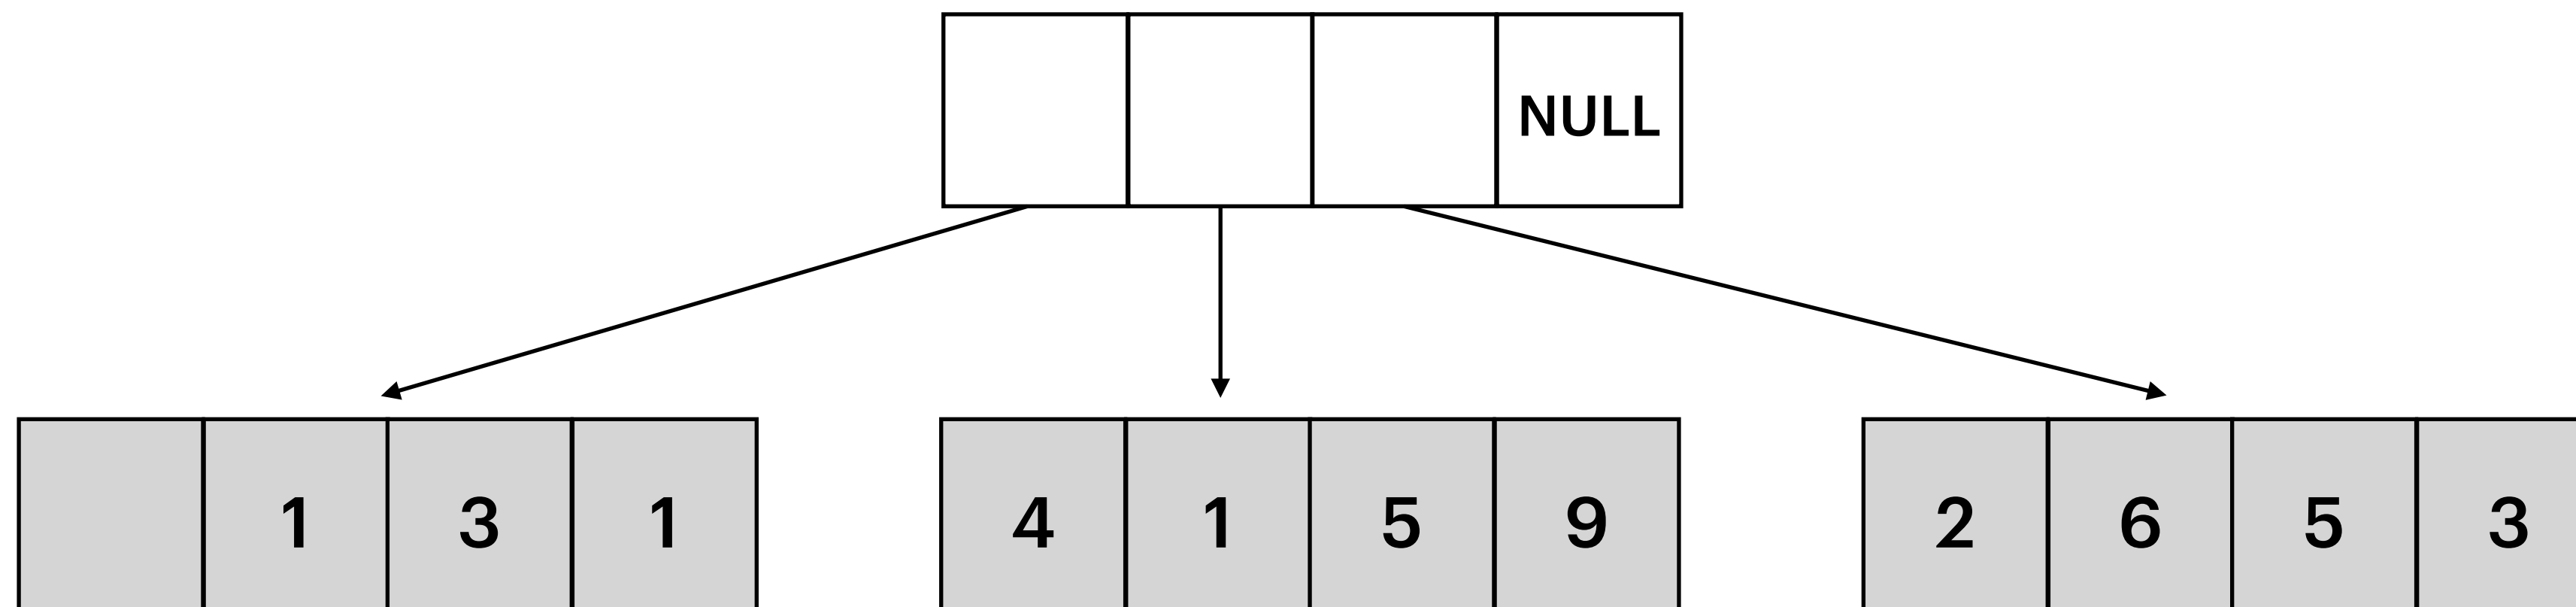
# How is a deque actually implemented?

- So how do we **push\_front(1)**?



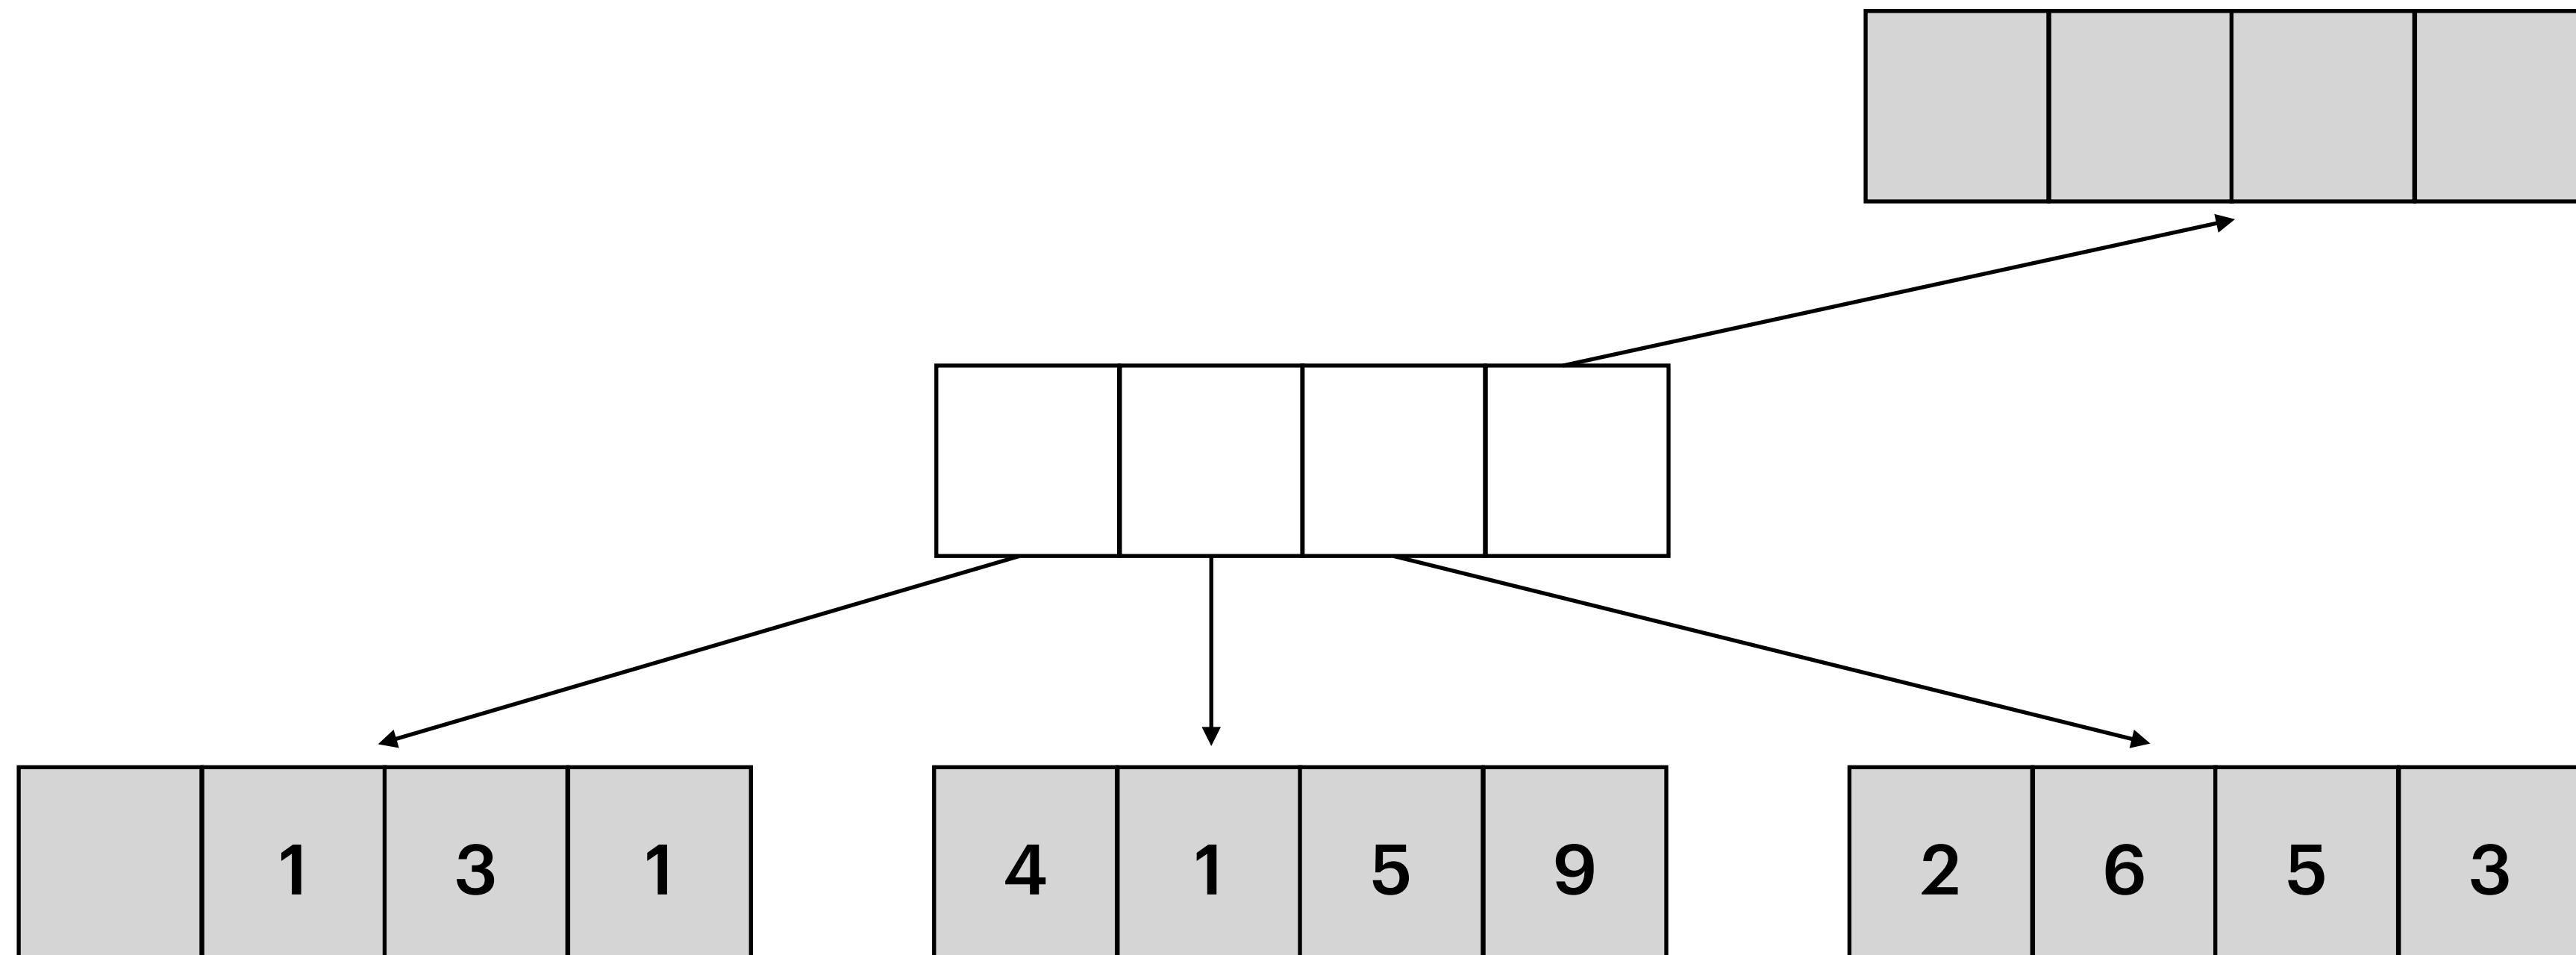
# How is a deque actually implemented?

- Now, how do we **push\_back(7)**? Haven't we run out of space in the last array?



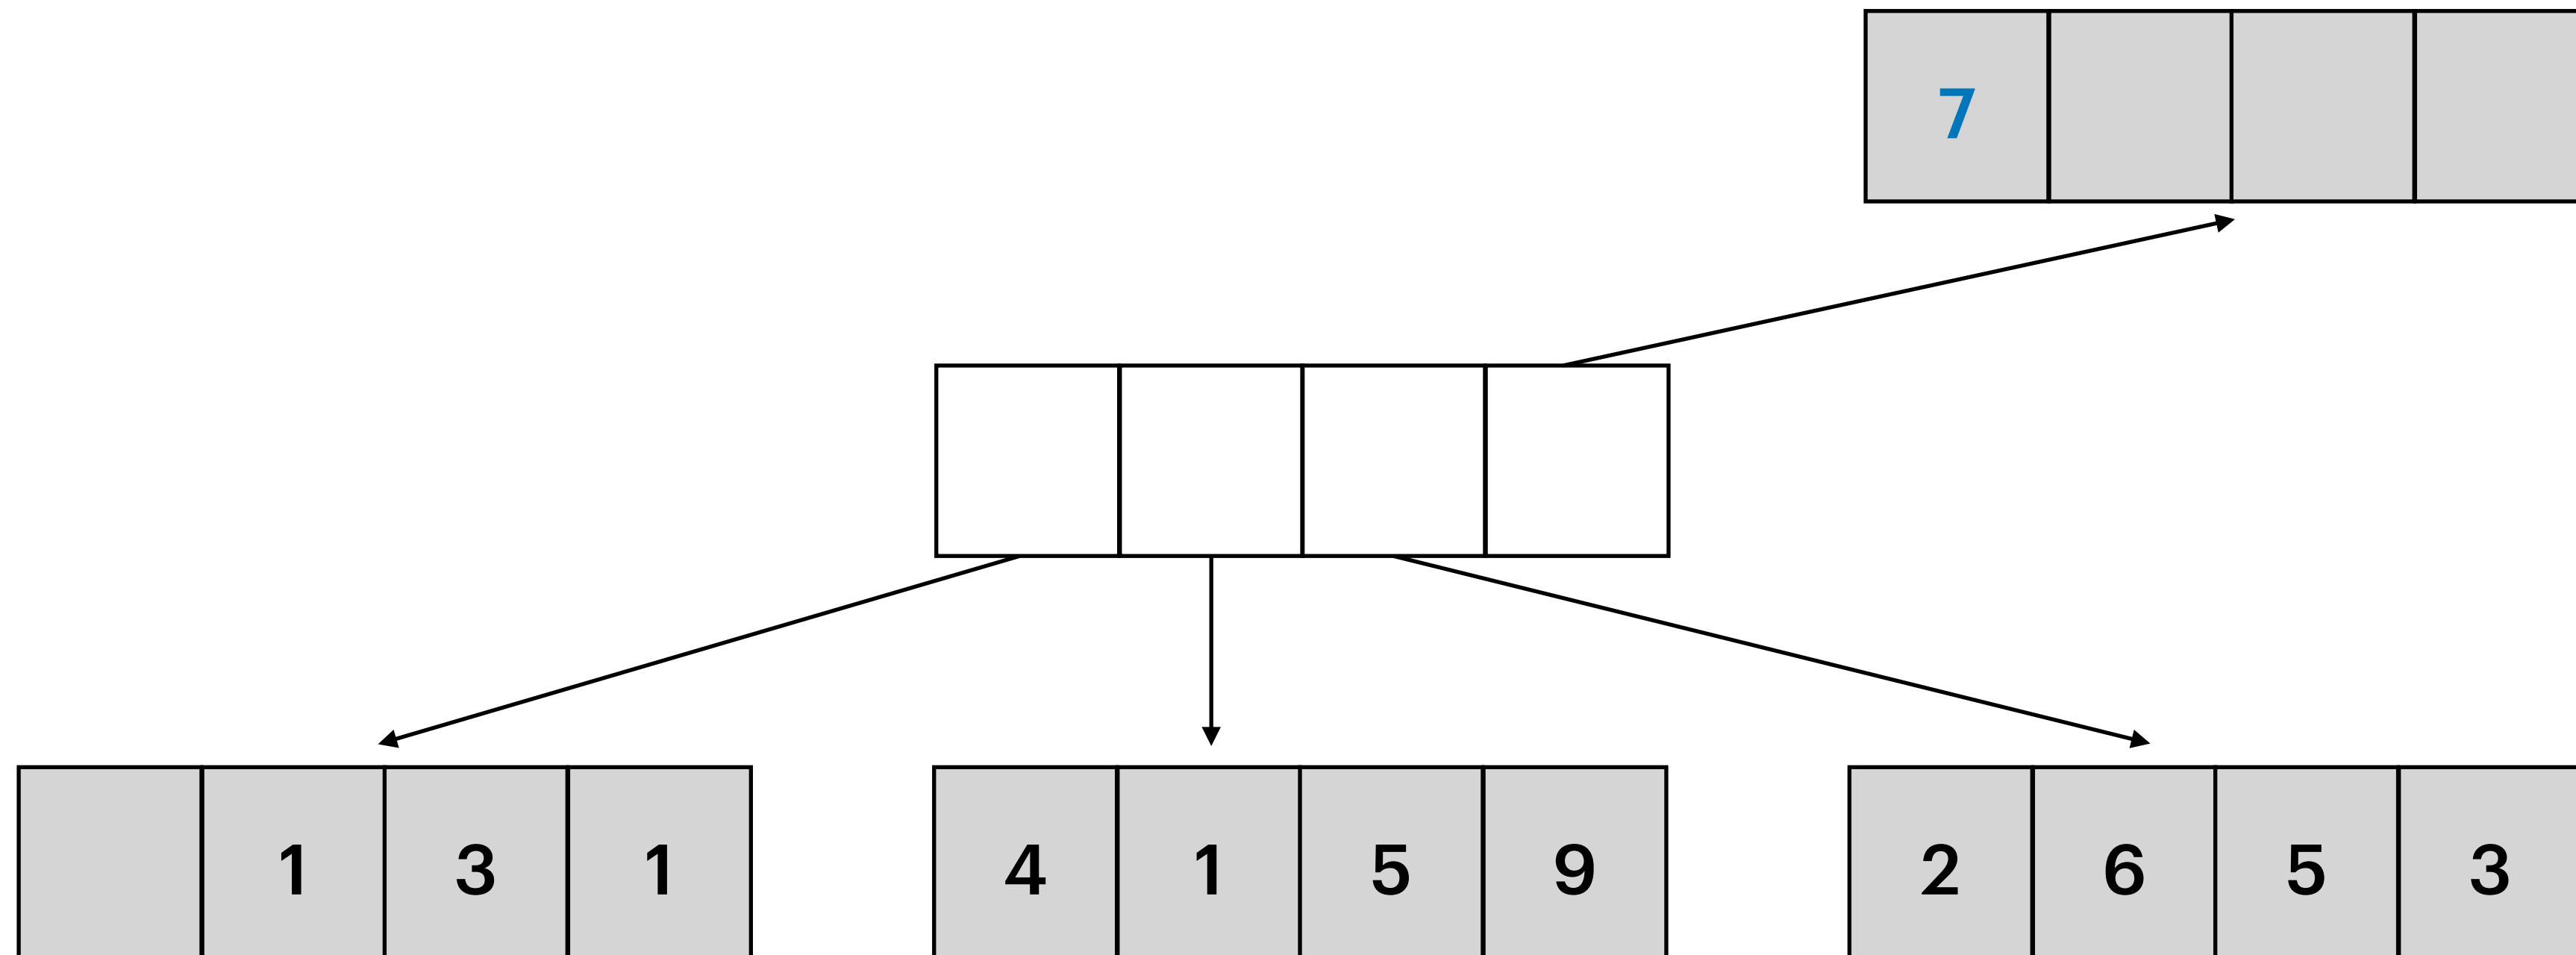
# How is a deque actually implemented?

- Now, how do we **push\_back(7)**? Haven't we run out of space in the last array?



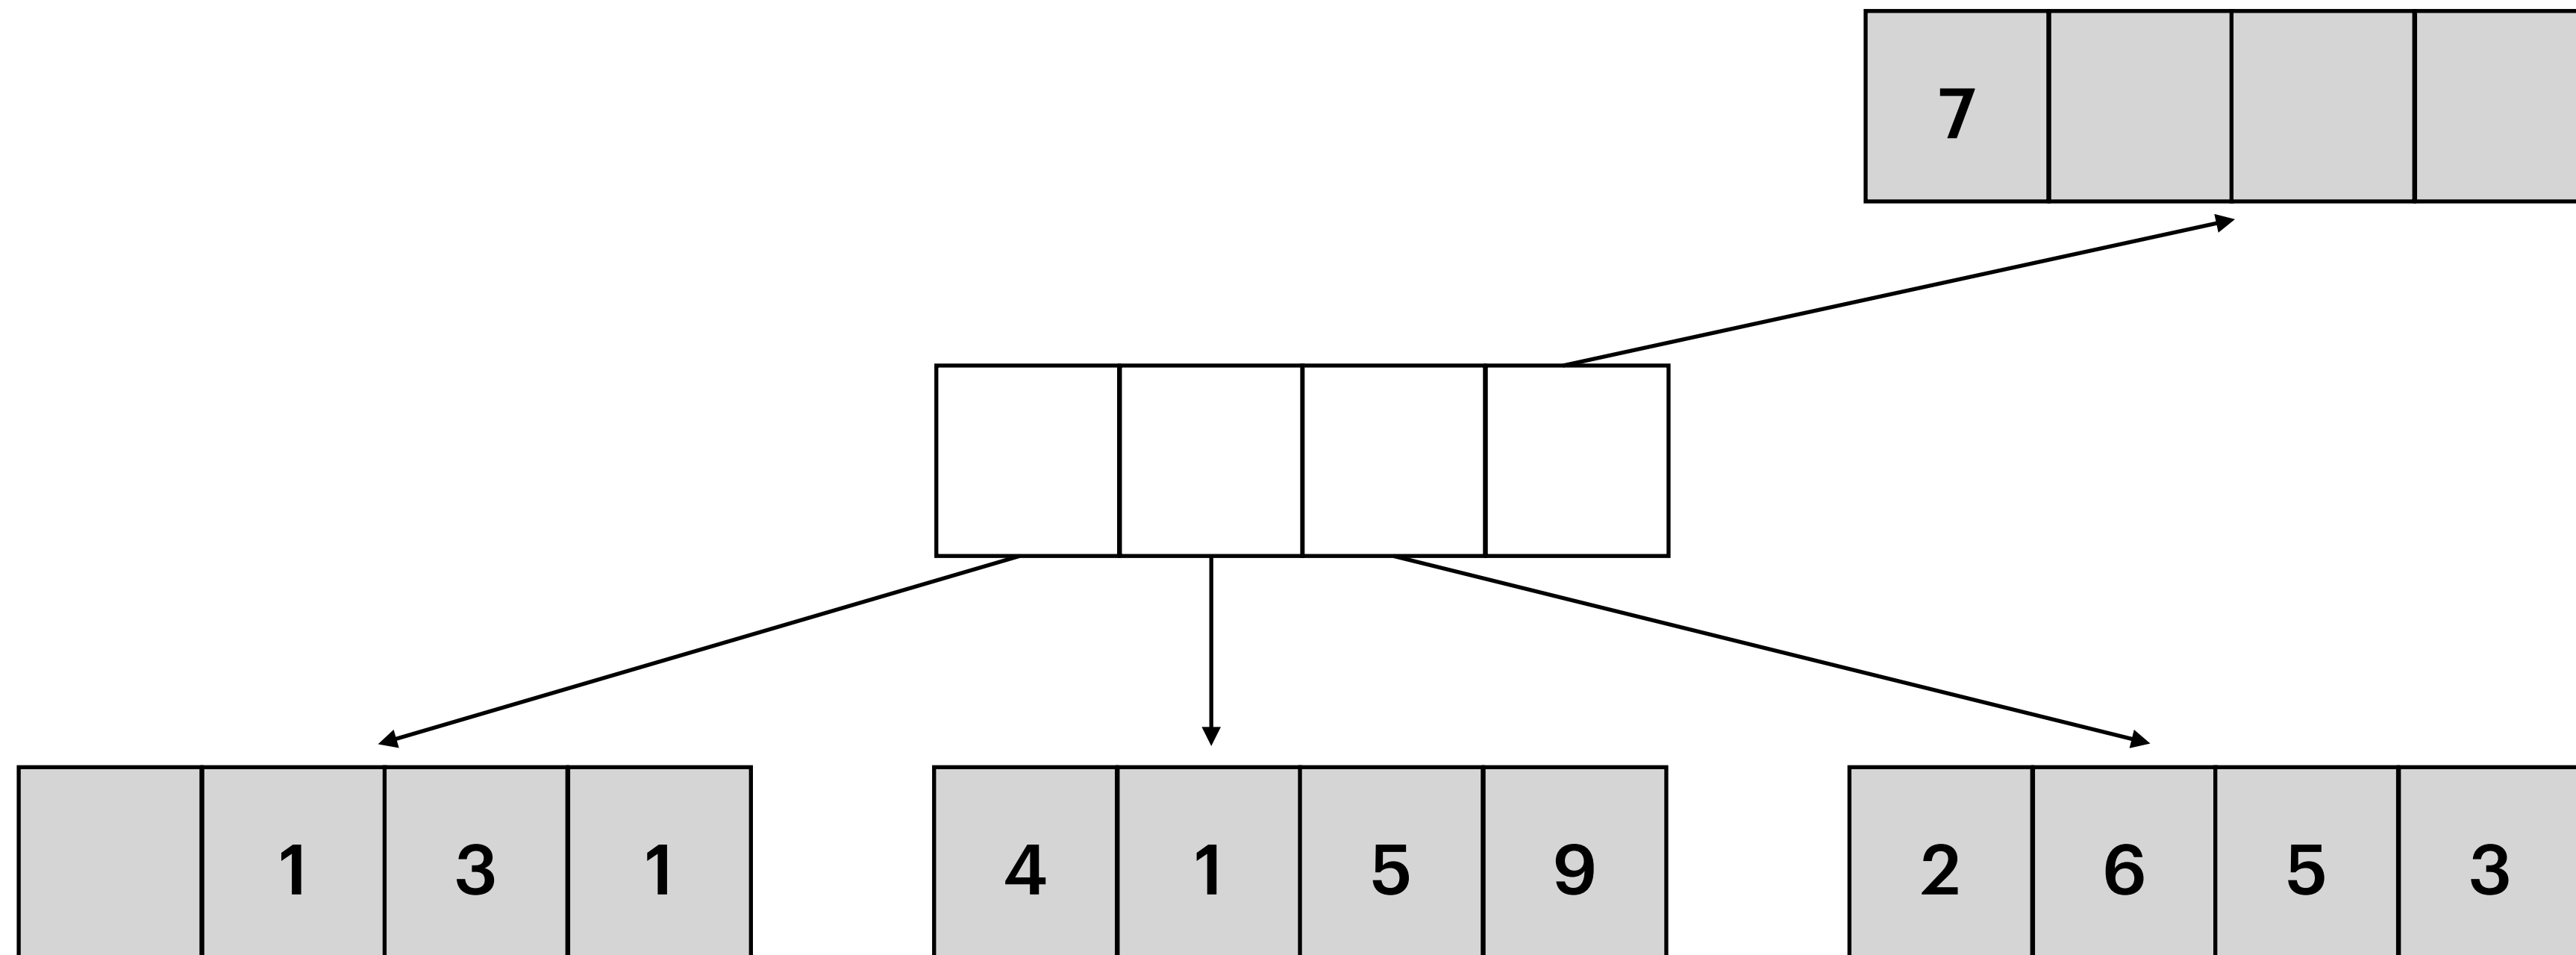
# How is a deque actually implemented?

- Now, how do we **push\_back(7)**? Haven't we run out of space in the last array?




# How is a deque actually implemented?

- Lastly, how can we **push\_front(8)** and then **push\_front(0)**?



# How is a deque actually implemented?

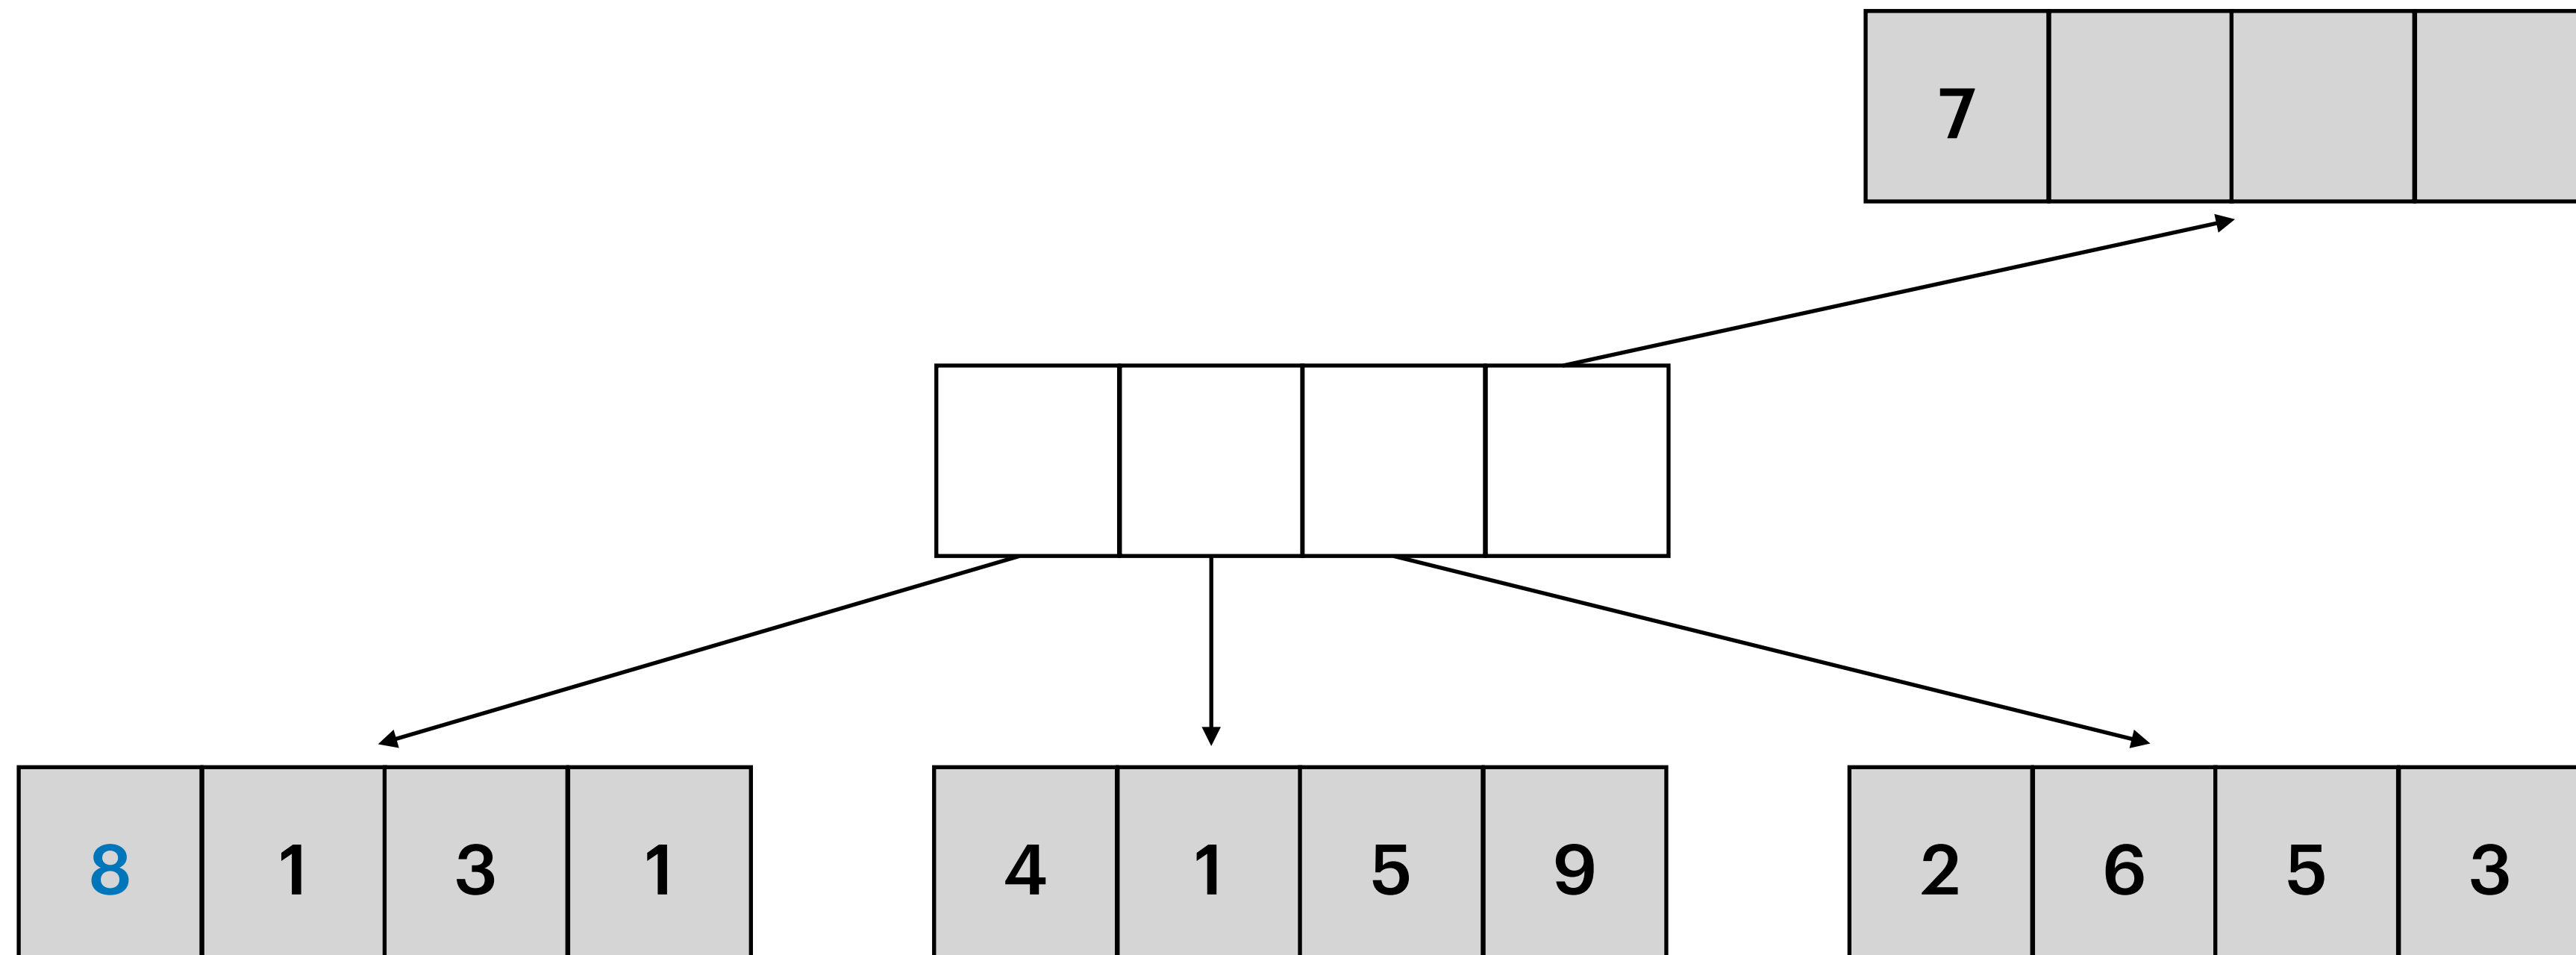
- Lastly, how can we **push\_front(8)** and then **push\_front(0)**?



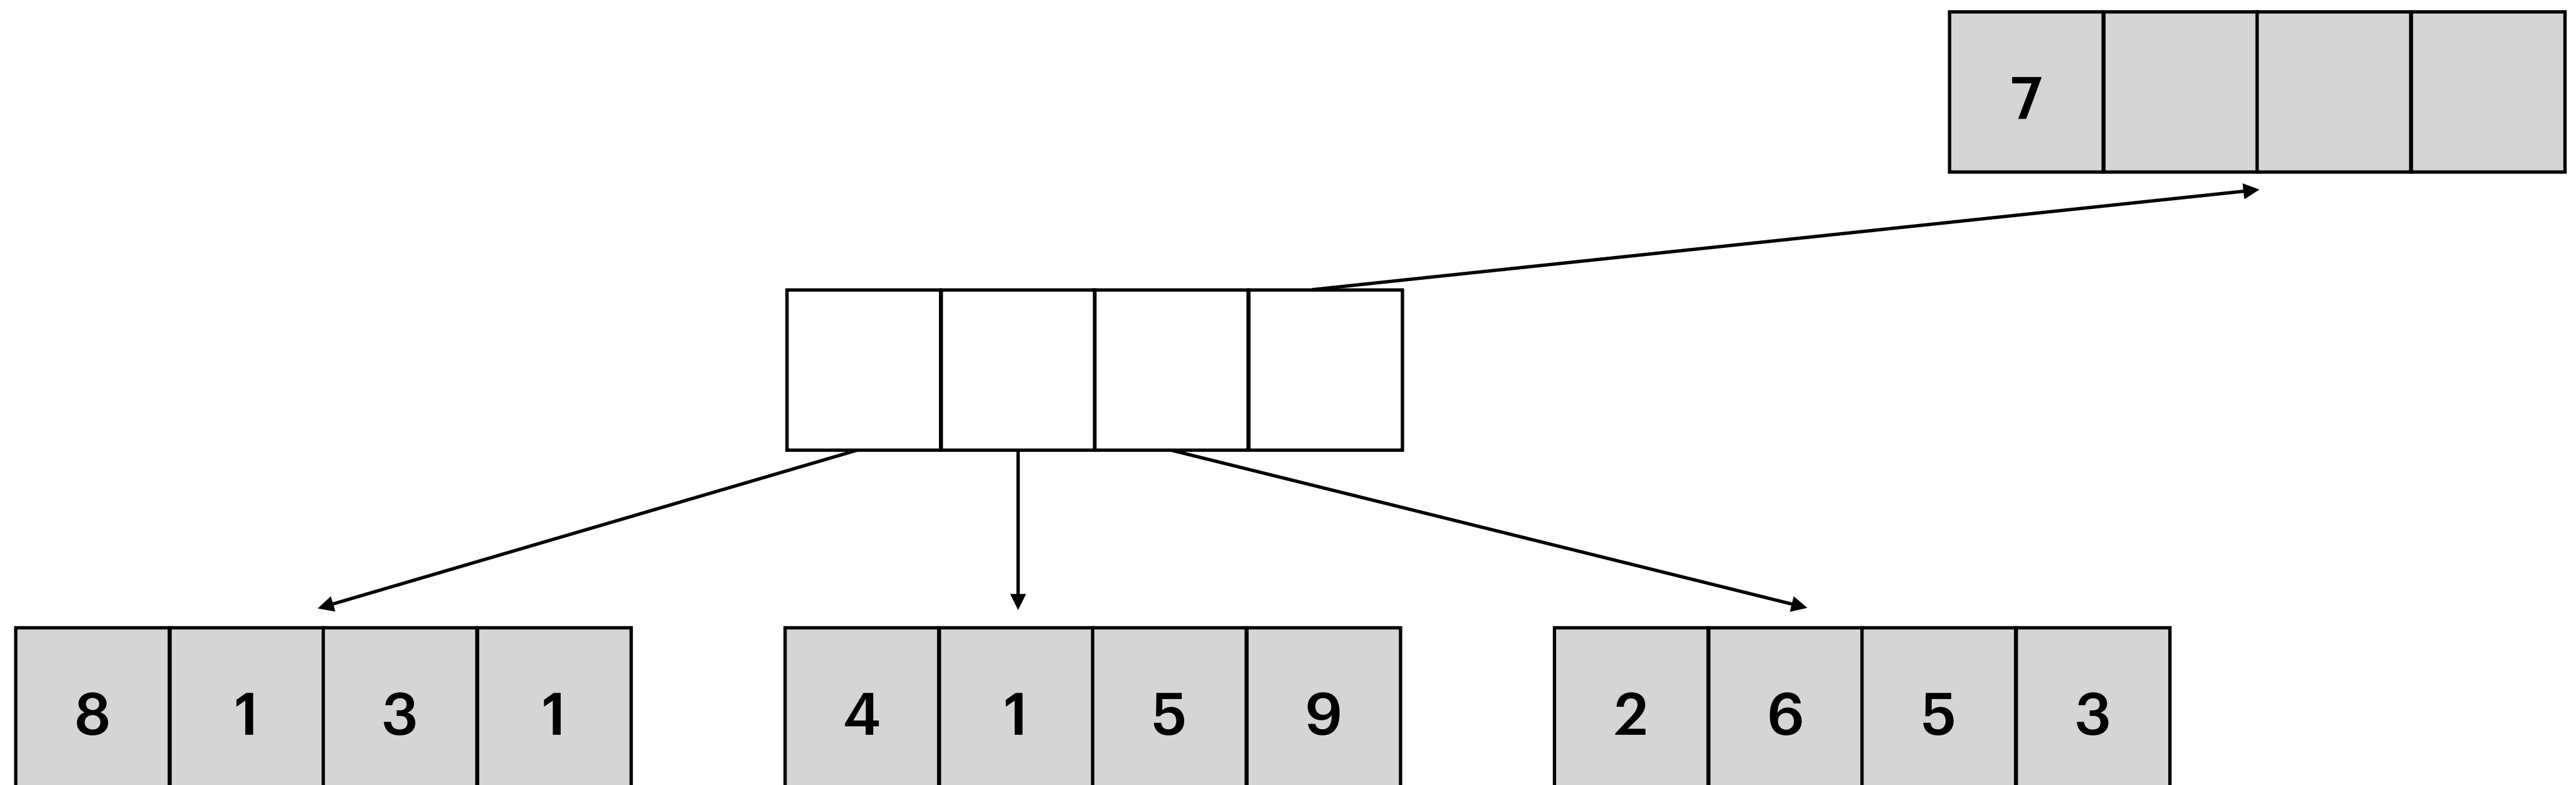


# How is a deque actually implemented?

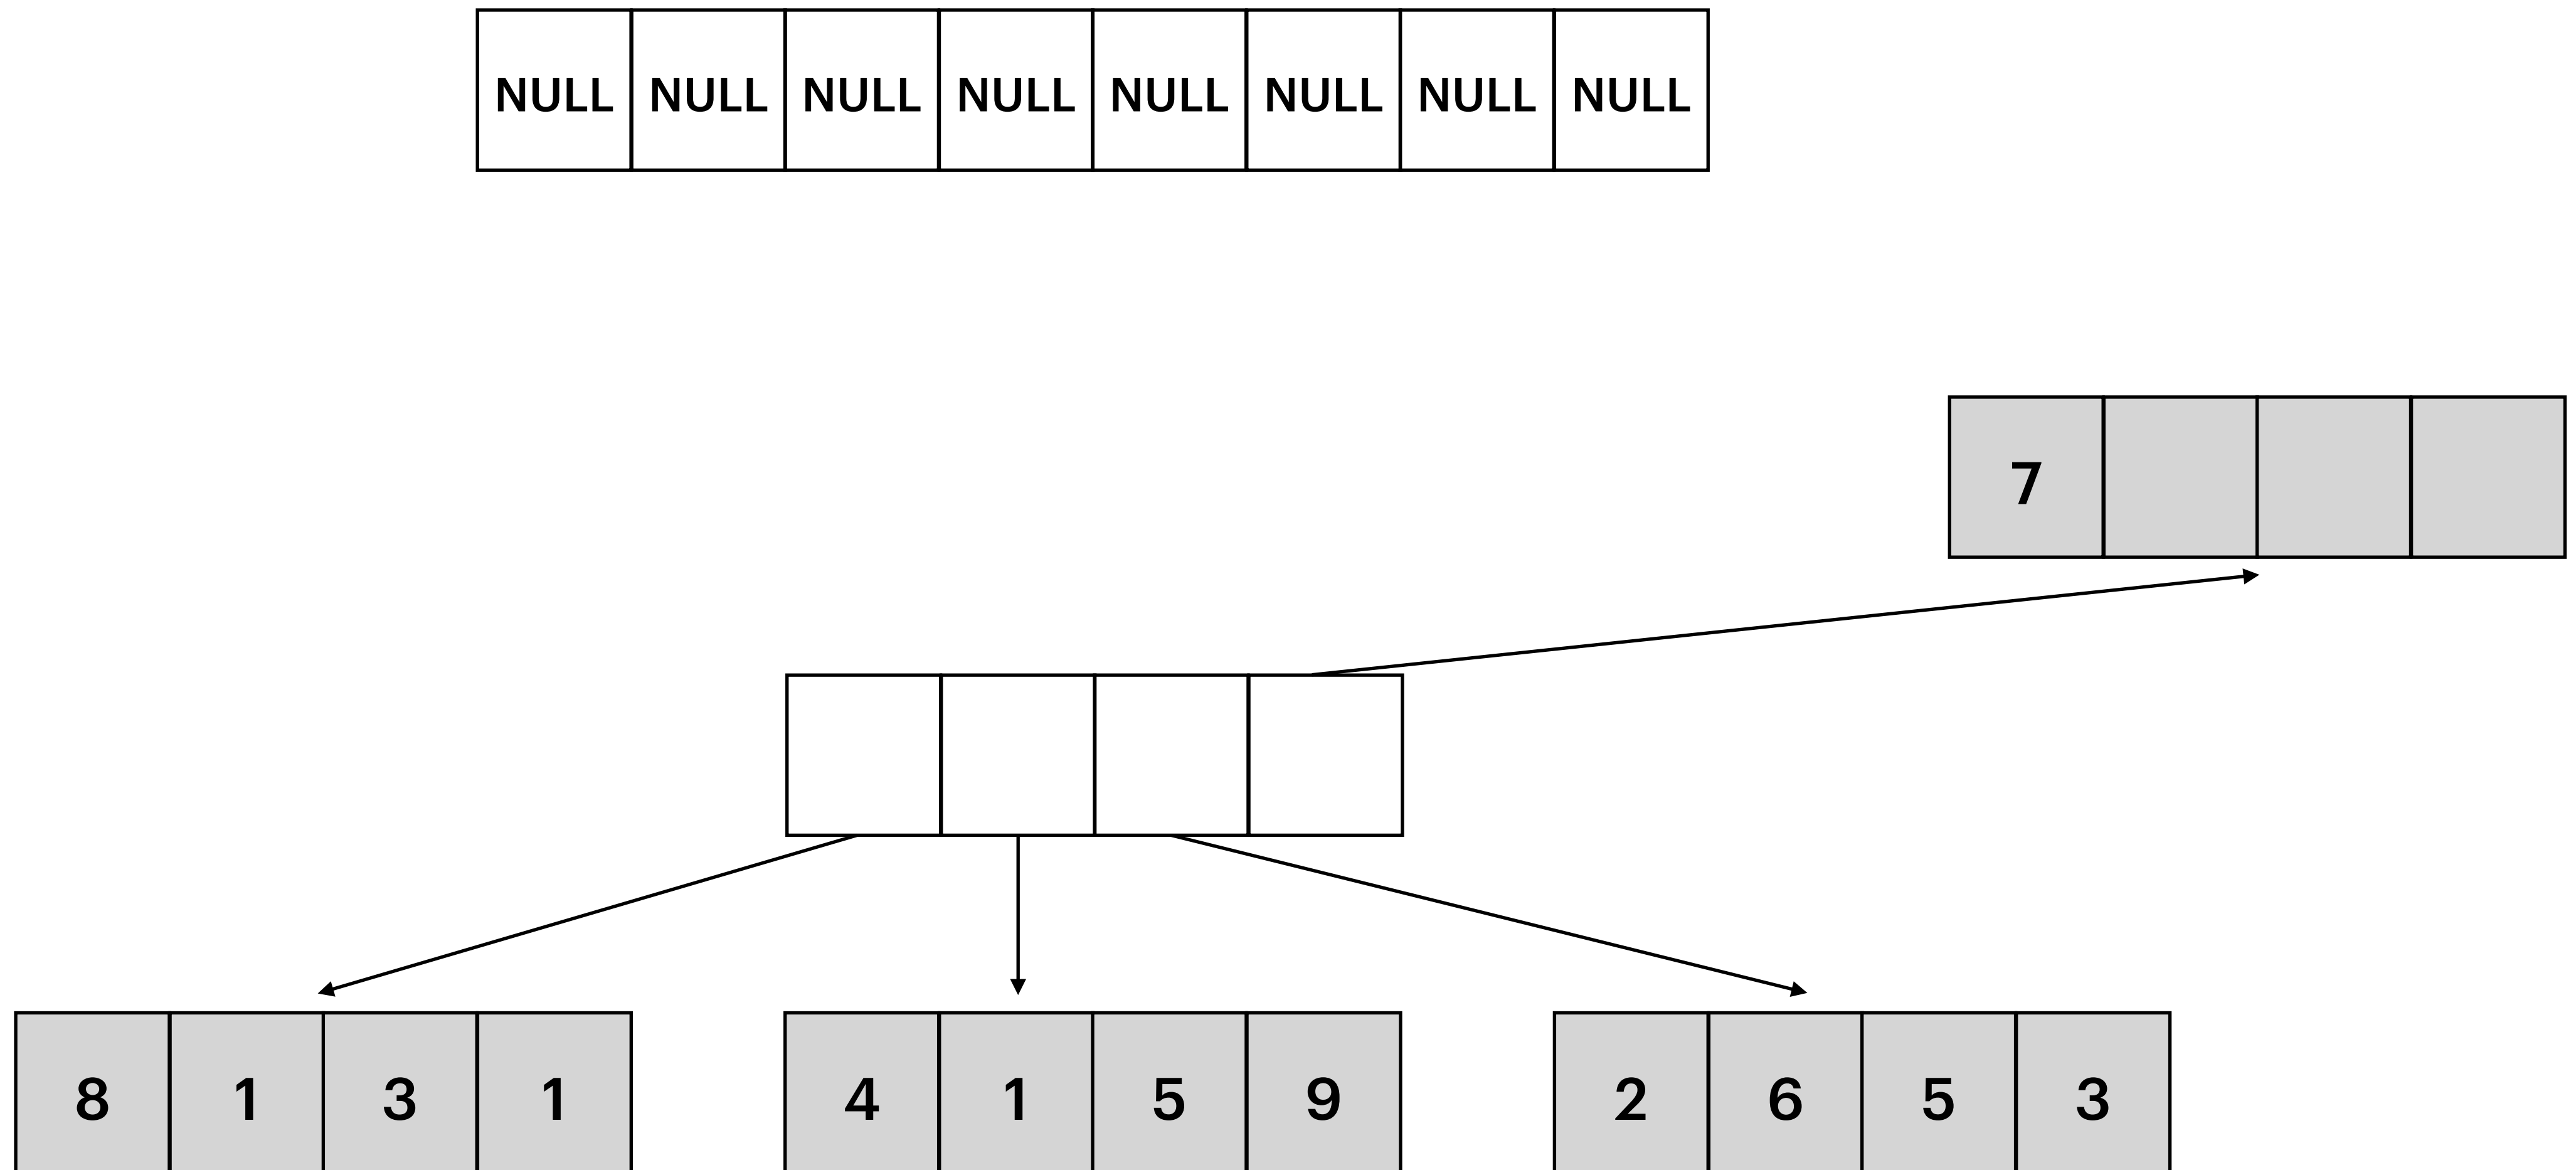
- Lastly, how can we **push\_front(8)** and then **push\_front(0)**?



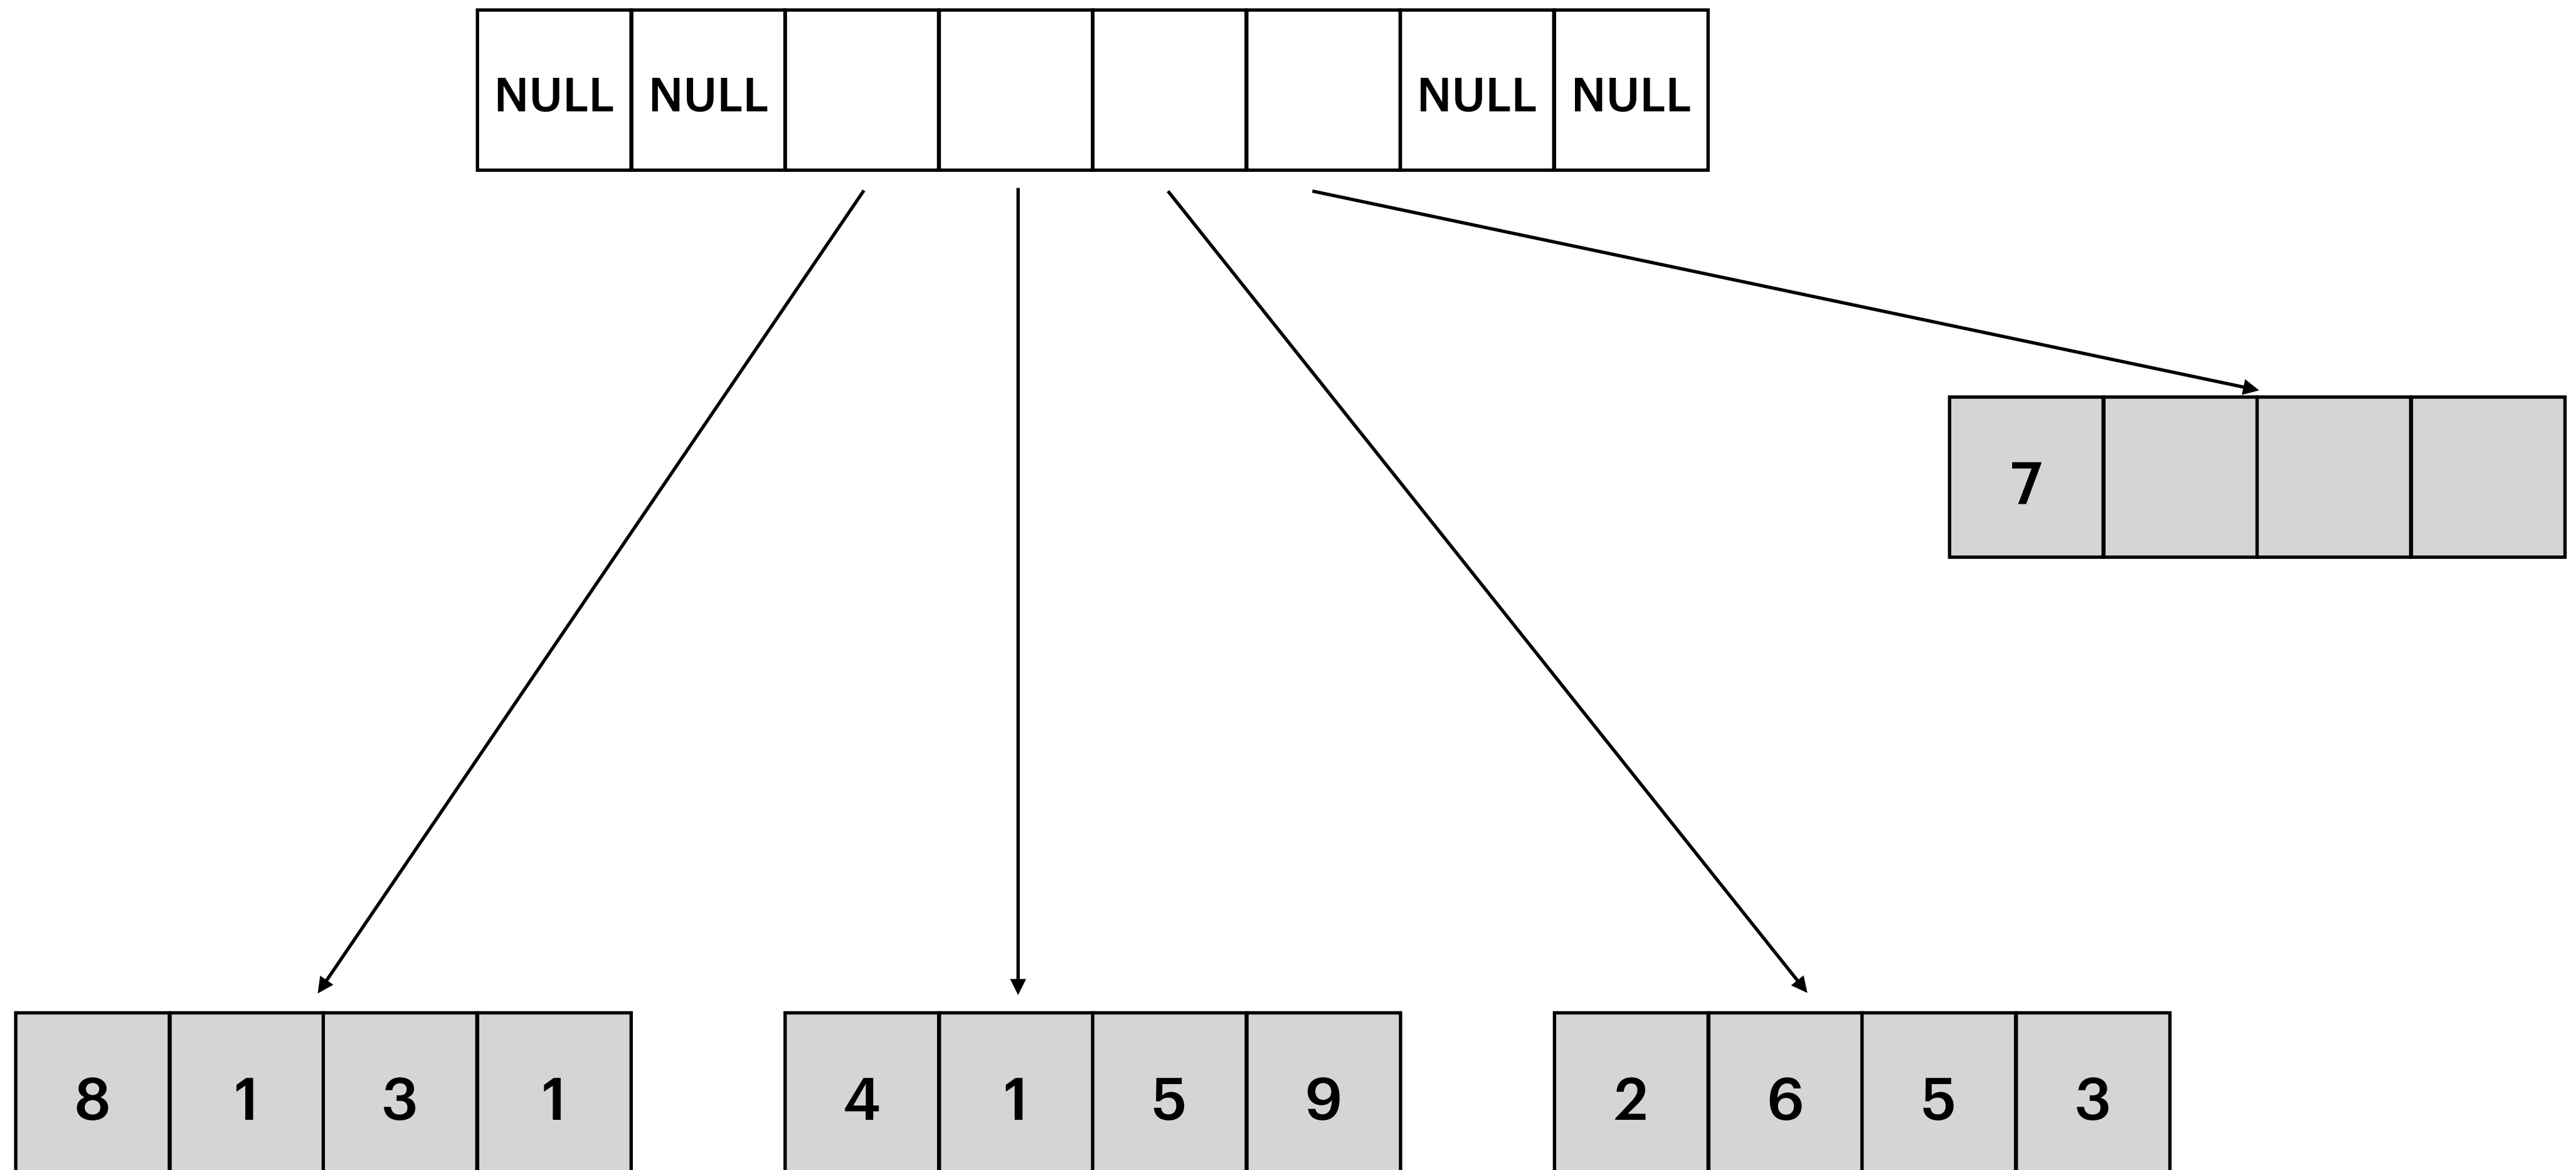
# How is a deque actually implemented?



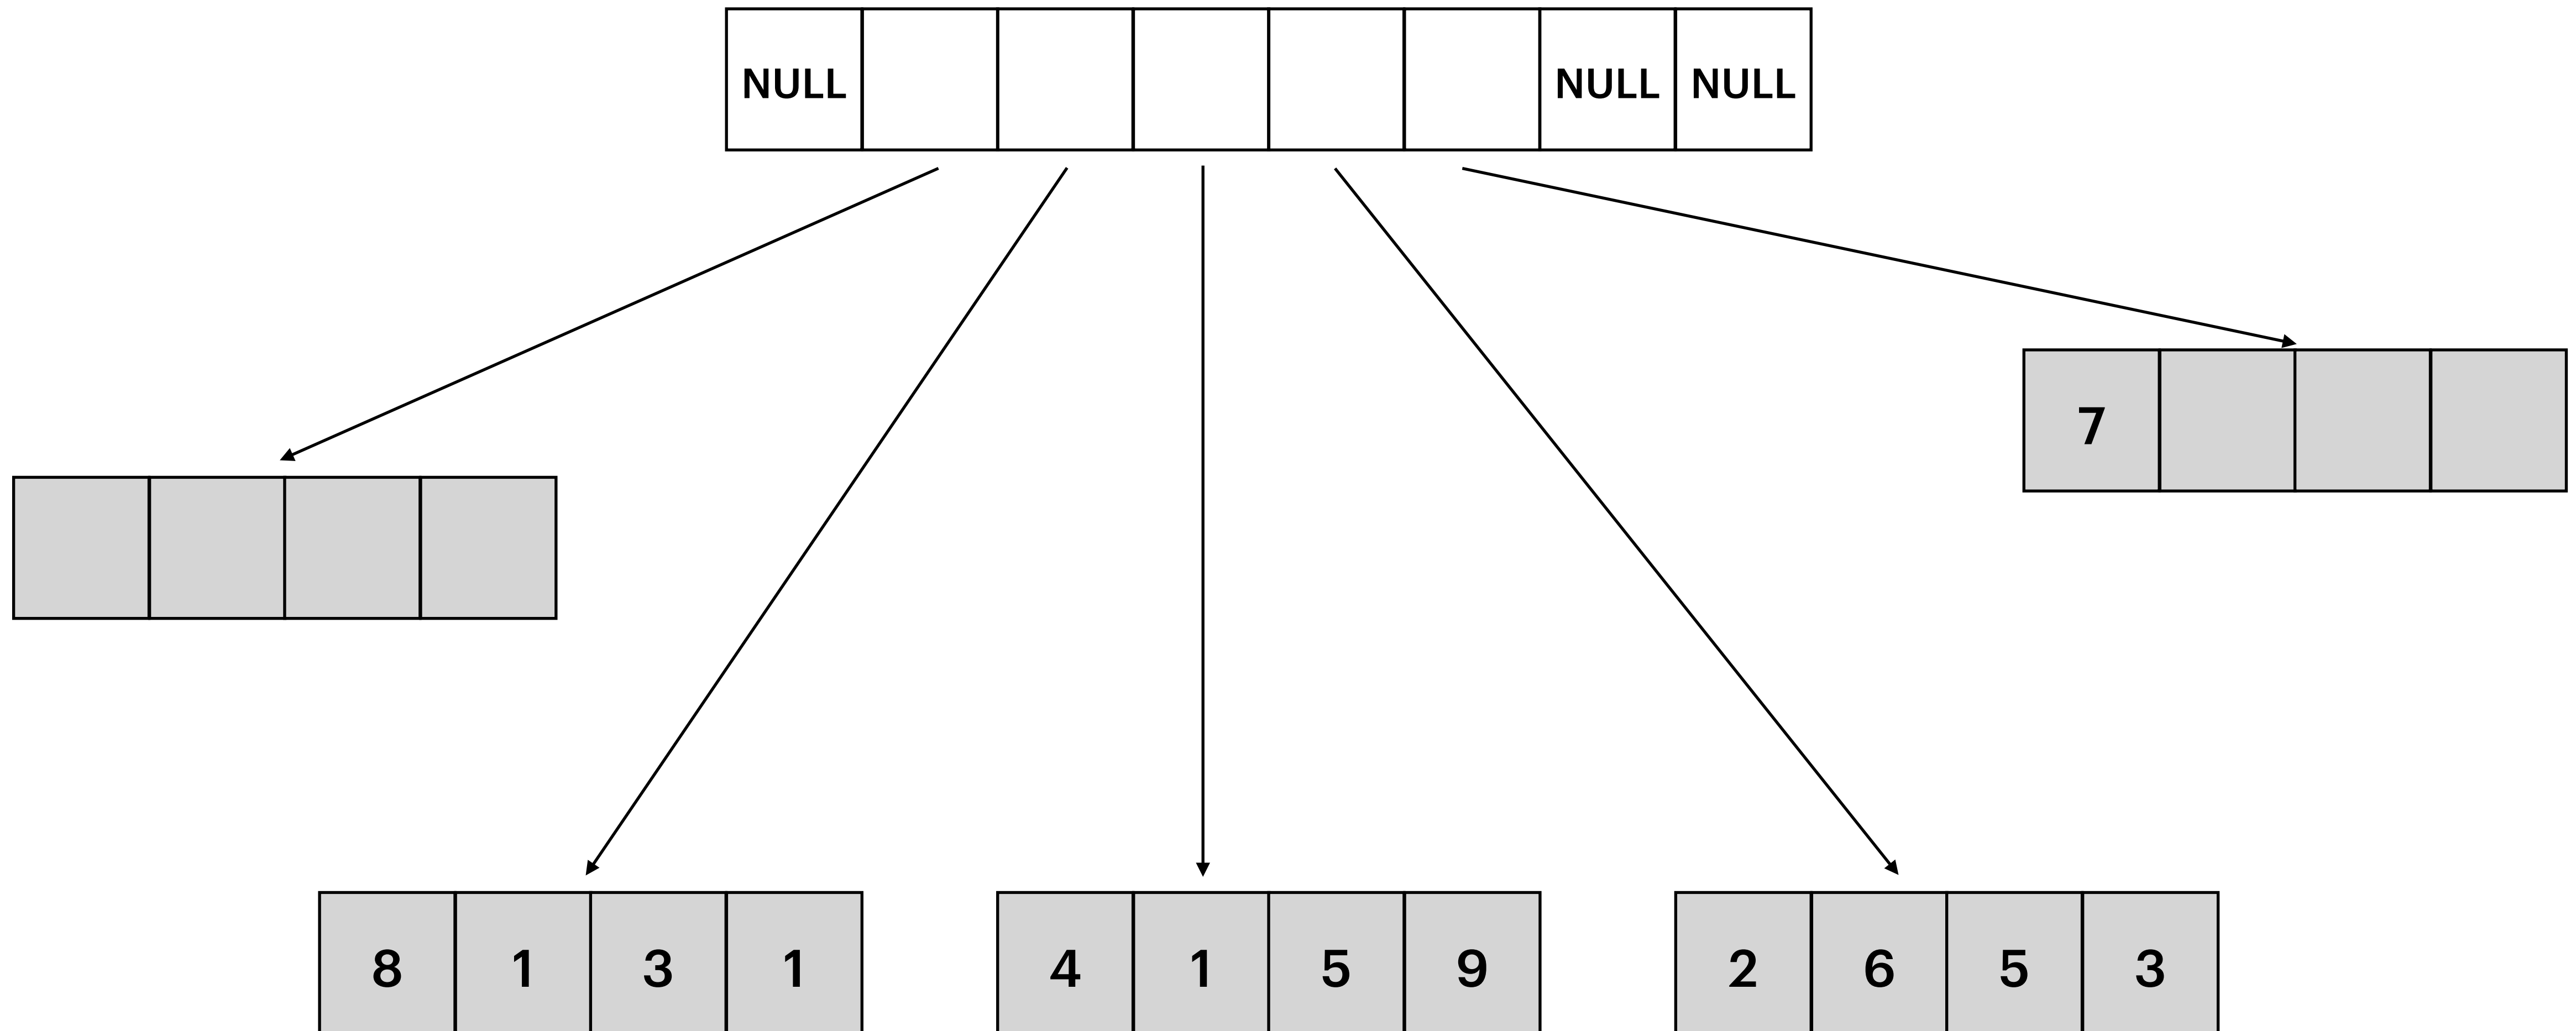
# How is a deque actually implemented?



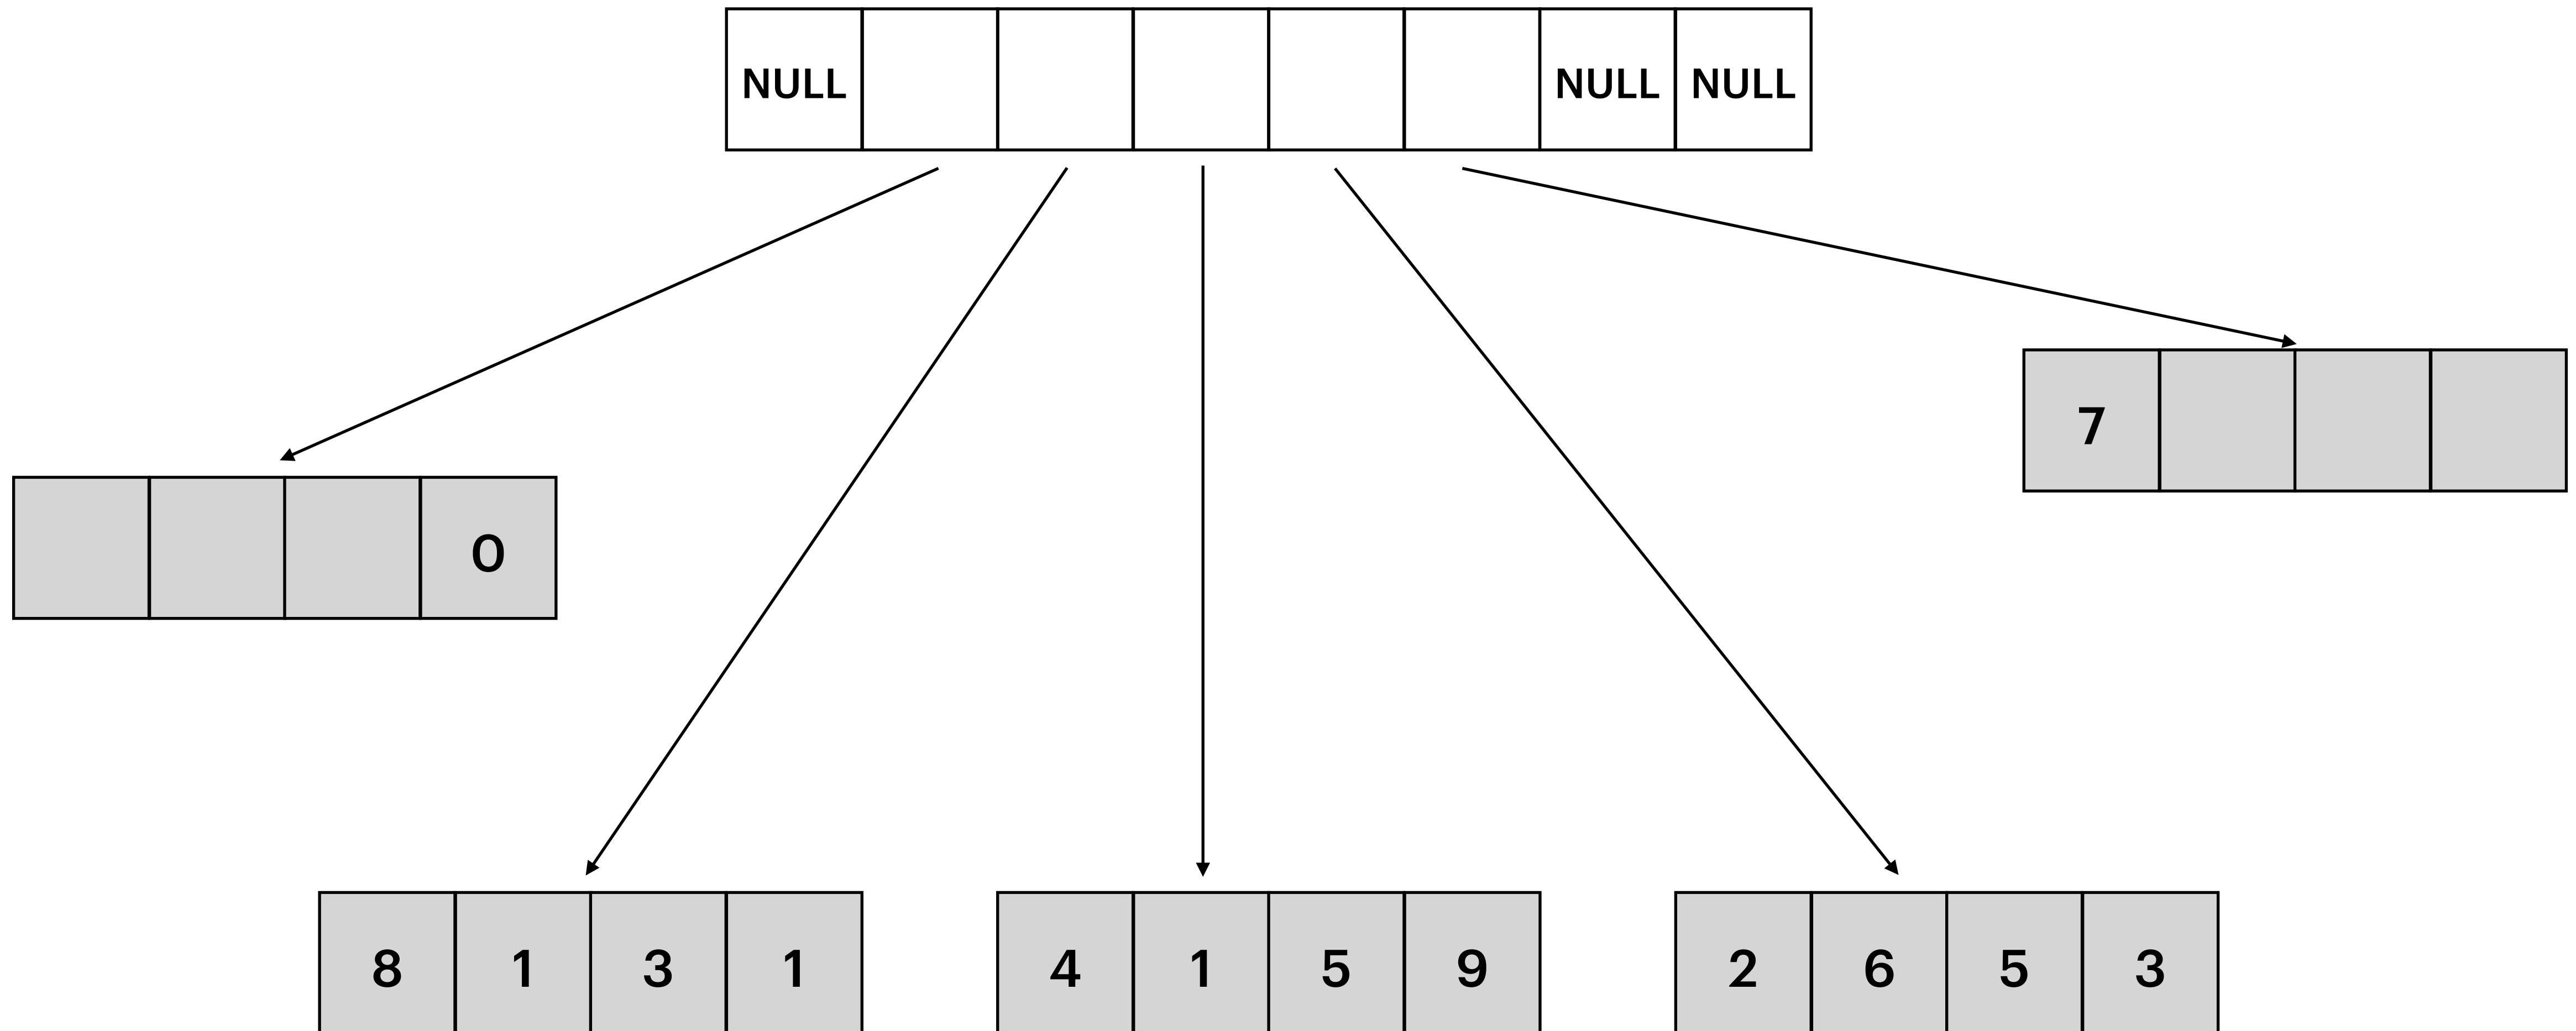
# How is a deque actually implemented?



# How is a deque actually implemented?



# How is a deque actually implemented?



# How is a list actually implemented?

- Recap: a list provides fast insertion anywhere, but no random (indexed) access

# How is a list actually implemented?

- Recap: a list provides fast insertion anywhere, but no random (indexed) access

```
std::list<int> list{5, 6};           // {5, 6}
list.push_front(3);                  // {3, 5, 6}
list.pop_back();                     // {3, 5}
```

Usually a doubly linked list. There's also a `forward_list` that's singly linked. Linked lists will be covered at the end of 106B, so don't fret if this footnote is unfamiliar to you!



# When to use which sequence container?

What you want to do	<code>std::vector</code>	<code>std::deque</code>	<code>std::list</code>
Insert/remove in the front	Slow	Fast	Fast
<b>Insert/remove in the back</b>	Super Fast	Very Fast	Fast
<b>Indexed Access</b>	Super Fast	Fast	Impossible
Insert/remove in the middle	Slow	Fast	Very Fast
Memory usage	Low	High	High
Combining (splicing/joining)	Slow	Very Slow	Fast
Stability* (iterators/concurrency)	Bad	Very Bad	Good

These two  
are the most  
common!

Don't worry if you don't know what stability means!  
It's a fairly advanced concept that you don't need to  
understand in order to grasp the core of this slide.

# Summary of Sequence Containers

**`std::vector`: use for almost everything**

`std::deque`: use if you are frequently inserting/removing at front


`std::list`: use very rarely, and only if you need to split/join multiple lists

# Container Adaptors

What is a container adaptor?  
`std::stack` and `std::queue`

# Types of containers

- All containers can hold almost all elements.



<> vector

↕ deque

↓ list


() tuple




stack



queue



priority\_queue





(adding/removing elements from the front)

(adding elements from the front, removing from the back)

(adding elements with a priority, always removing the highest priority-element)


# What is a wrapper?

- A wrapper on an object changes how external users can interact with that object.




# What is a wrapper?

- The vault's owner has access to all possible ways to use the vault!
- Should a vault customer be able to do the same actions as the vault owner?




# What is a wrapper?

- Of course not! The vault teller limits your access to the vault.
- The teller is in charge of forwarding your requests to the actual vault.



# What is a wrapper?


- Of course not! The vault teller limits your access to the vault.
- The teller is in charge of forwarding your requests to the actual vault.





# Container adaptors are wrappers in C++!


- Container adaptors provide a different interface for sequence containers.
- You can choose what the underlying container is!
- For instance, let's choose a deque as our underlying container, and let's implement a queue!



# Container adaptors are wrappers in C++!

- Container adaptors provide a different interface for sequence containers.
- You can choose what the underlying container is!
- For instance, let's choose a deque as our underlying container, and let's implement a queue!


`std::queue`



# Container adaptors are wrappers in C++!

- Container adaptors provide a different interface for sequence containers.
- You can choose what the underlying container is!
- For instance, let's choose a deque as our underlying container, and let's implement a queue!

`std::queue`



# std::stack and std::queue

## std::queue

---

Defined in header `<queue>`

```
template<
    class T,
    class Container = std::deque<T>
> class queue;
```

---

The `std::queue` class is a **container adapter** that gives the programmer the functionality of a queue - specifically, a FIFO (first-in, first-out) data structure.

The class template acts as a **wrapper** to the underlying container - **only a specific set of functions is provided**. The queue pushes the elements on the back of the underlying container and pops them from the front.

---

## std::stack

---

Defined in header `<stack>`

```
template<
    class T,
    class Container = std::deque<T>
> class stack;
```

---

The `std::stack` class is a container adapter that gives the programmer the functionality of a stack - specifically, a LIFO (last-in, first-out) data structure.

The class template acts as a wrapper to the underlying container - only a specific set of functions is provided. The stack pushes and pops the element from the back of the underlying container, known as the top of the stack.

# Concrete examples with `std::queue`

## `std::queue`

Defined in header `<queue>`

```
template<
    class T,
    class Container = std::deque<T>
> class queue;
```

The `std::queue` class is a container adapter that gives the programmer the functionality of a queue - specifically, a FIFO (first-in, first-out) data structure.

The class template acts as a **wrapper** to the underlying container - **only a specific set of functions is provided**. The queue **pushes the elements** on the back of the underlying container and pops them from the front.

```
std::queue<int> stack_deque;           // Container = std::deque

std::queue<int, std::list<int>> stack_list; // Container = std::list

std::queue<int, std::vector<int>> stack_vector; // Container = std::vector?
```



# Concrete examples with `std::queue`

## `std::queue`

Defined in header `<queue>`

```
template<
    class T,
    class Container = std::deque<T>
> class queue;
```

The `std::queue` class is a container adapter that gives the programmer the functionality of a queue - specifically, a FIFO (first-in, first-out) data structure.

The class template acts as a **wrapper** to the underlying container - **only a specific set of functions is provided**. The queue **pushes the elements** on the back of the underlying container and pops them from the front.

```
std::queue<int> stack_deque;           // Container = std::deque

std::queue<int, std::list<int>> stack_list; // Container = std::list

std::queue<int, std::vector<int>> stack_vector; // Container = std::vector
```

removing from the front of a vector is slow!

# Some member functions of `std::queue`

## Member functions

(constructor)	constructs the queue (public member function)
(destructor)	destructs the queue (public member function)
<b>operator=</b>	assigns values to the container adaptor (public member function)

## Element access

<b>front</b>	access the first element (public member function)
<b>back</b>	access the last element (public member function)

## Capacity


<b>empty</b>	checks whether the underlying container is empty (public member function)
<b>size</b>	returns the number of elements (public member function)

## Modifiers

<b>push</b>	inserts element at the end (public member function)
<b>emplace</b> (C++11)	constructs element in-place at the end (public member function)
<b>pop</b>	removes the first element (public member function)
<b>swap</b> (C++11)	swaps the contents (public member function)

# Types of containers

- All containers can hold almost all elements.



<> vector

↕ deque

↓ list


() tuple




stack



queue



priority\_queue



(adding/removing elements from the front)

(adding elements from the front, removing from the back)

(adding elements with a priority, always removing the highest priority-element)

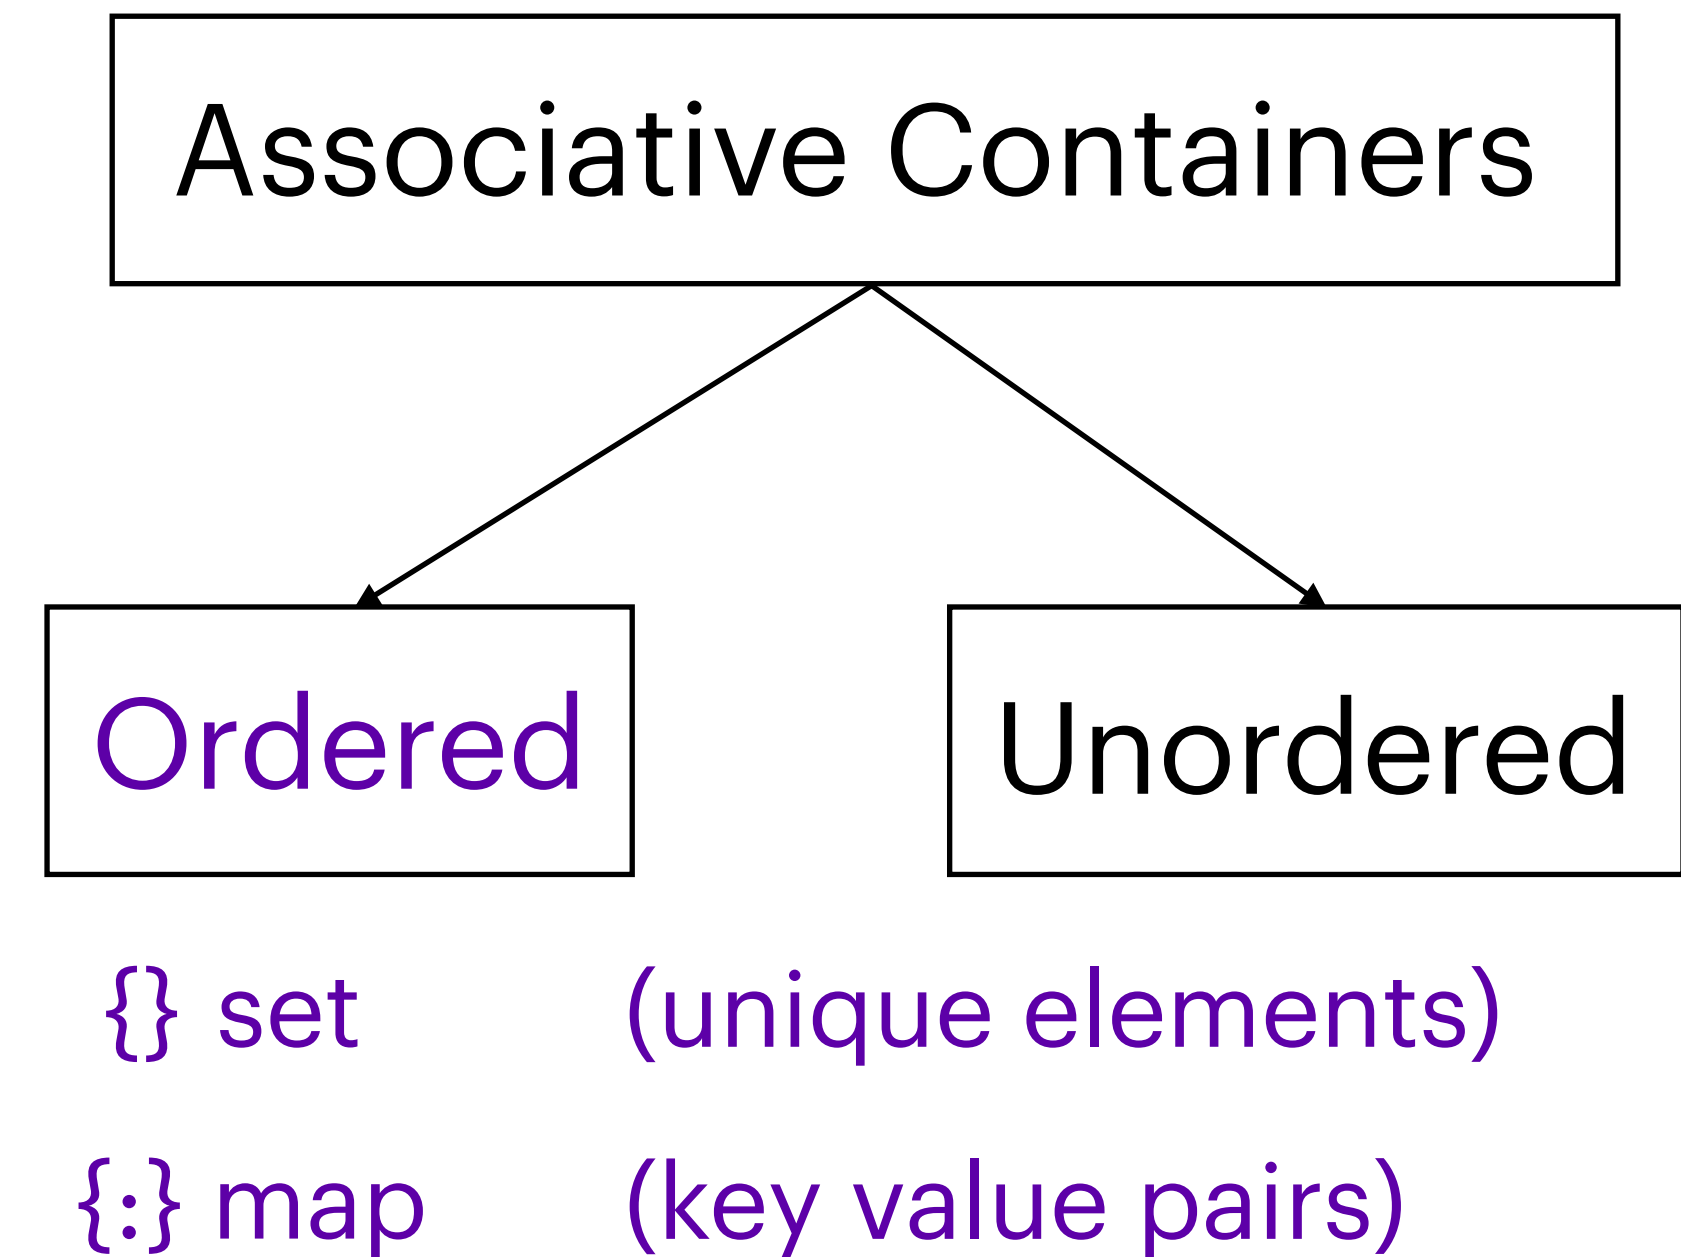
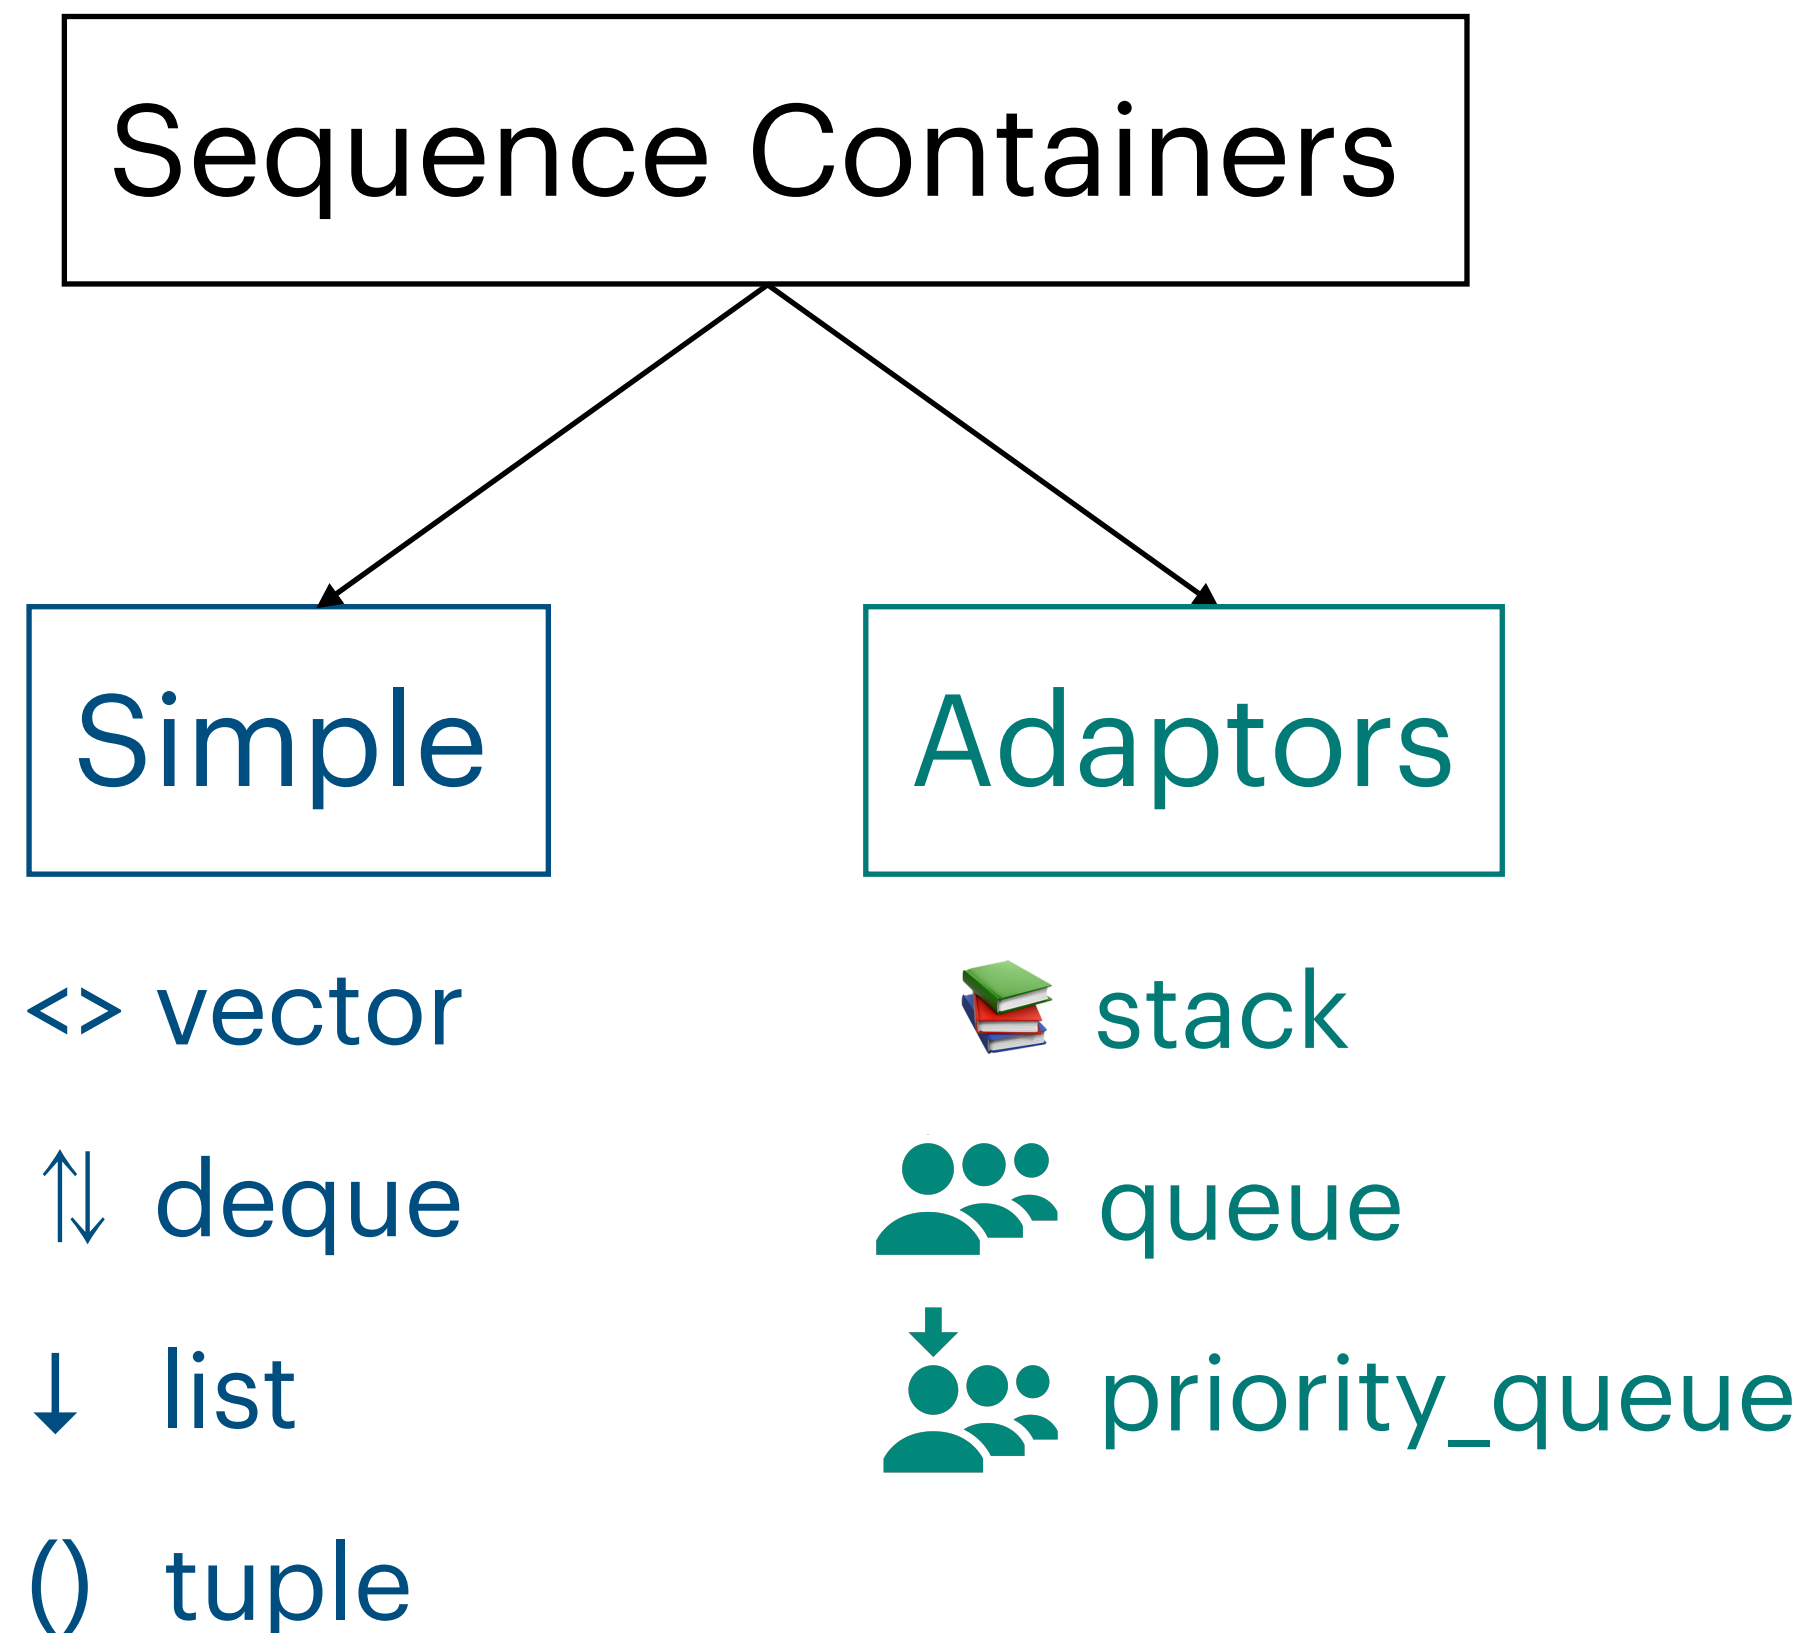


# Associative Containers

`std::set` functions  
`std::map` functions and auto-insertion  
type requirements

# Types of containers

- All containers can hold almost all elements.



# Live Demo!

Let's compare the Stanford's Map/Set and STL's map/set:  
(QT Creator Project)

# Stanford “Set” vs STL “set”

What you want to do	Stanford Set<int>	std::set<int>
Create an empty set	Set<int> s;	std::set<int> s;
Add a value <b>k</b> to the set	s.add(k);	s.insert(k);
Remove value <b>k</b> from the set	s.remove(k);	s.erase(k);
Check if a value <b>k</b> is in the set	if (s.contains(k)) ...	if (s.count(k)) ...
Check if vector is empty	if (vec.isEmpty()) ...	if (vec.empty()) ...

# Stanford “Map” vs STL “map”

What you want to do	Stanford Map<int, char>	std::map<int, char>
Create an empty map	Map<int, char> m;	std::map<int, char> m;
Add key k with value v into the map	m.put(k, v); m[k] = v;	m.insert({k, v}); m[k] = v;
Remove key k from the map	m.remove(k);	m.erase(k);
Check if key k is in the map	if (m.containsKey(k)) ...	if (m.count(k)) ...
Check if the map is empty	if (m.isEmpty()) ...	if (m.empty()) ...
Retrieve or overwrite value associated with key k ( <b>error</b> if key isn't in map)	Impossible (but does auto-insert)	char c = m.at(k); m.at(k) = v;
Retrieve or overwrite value associated with key k ( <b>auto-insert</b> if key isn't in map)	char c = m[k]; m[k] = v;	char c = m[k]; m[k] = v;

# STL maps actually store pairs!

Every `std::map<K, V>` is actually backed by:

`std::pair<const K, V>`

- Why do pairs make sense here?
  - Why not just tuples?
- Why is it `const K` instead of just `K`?

hint: `std::pair`'s are just two-element tuples!

# Iterating through maps and sets

- Exactly the same as CS106B!
- Because maps are implemented with **std::pair**, you can use structured binding on them!

```
std::set<...> s;  
std::map<..., ...> m;  
  
for (const auto& element : s) {  
    // do stuff with element  
}  
  
for (const auto& [key, value] : m) {  
    // do stuff with key and value  
}
```

## Both Stanford and STL sets+maps require comparison operator!

- By default, the type (for sets) or key's type (for maps) must have a comparison operator (<) defined.

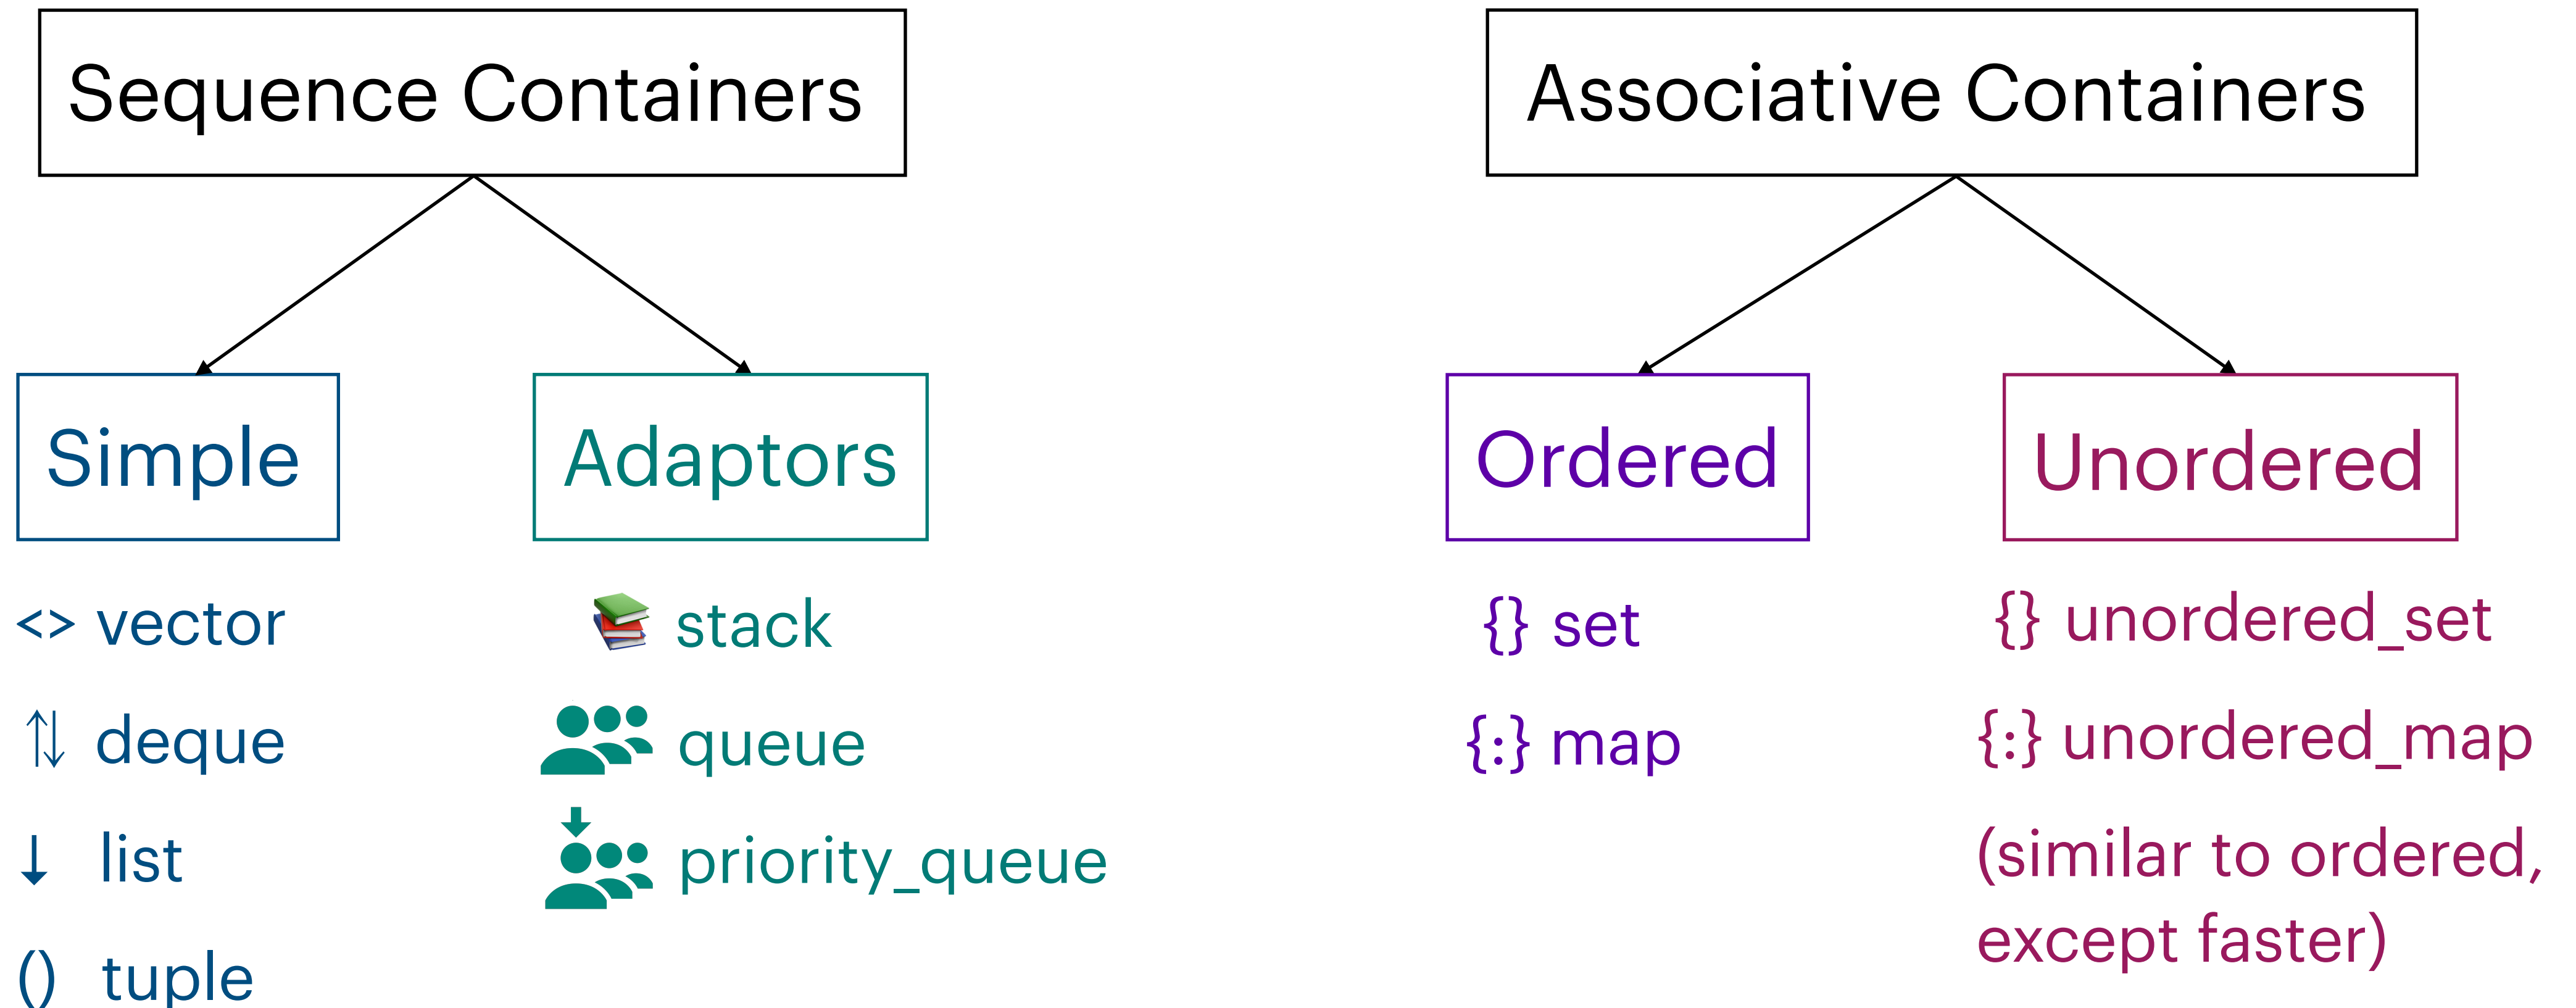
```
std::set<int> set1; // A OK! ints are comparable with <
std::set<std::ifstream> set2; // not ok. how do we compare ifstreams with < ?

std::map<int, int> map1; // A OK! ints are comparable with <
std::set<std::ifstream, int> map2; // not ok. how do we compare ifstreams with < ?
```



# Types of containers

- All containers can hold almost all elements.



# unordered\_map and unordered\_set

- Each STL **set/map** comes with an **unordered** sibling. They're almost the same, except:
  - Instead of a comparison operator, the set/map type must have a **hash function** defined for it.
    - Simple types, like **int**, **char**, **bool**, **double**, and even **std::string** are already supported!
    - Any containers/collections need you to provide a hash function to use them.
  - unordered\_map/unordered\_set are generally faster than map/set.

# That's a lot! Any broad tips for choosing one?

- What do they all have in common?
  - You can copy all of them!
  - You can check if two have the same elements in the same order.
  - You can get their size.
  - You can use iterators to access them (next lecture).

# That's a lot! Any broad tips for choosing one?

- How are they different?
  - Most containers can hold any data type.
  - Unordered associative containers (sets and maps) are tricky to get working if the element or key is another collection.

# Recap of STL Containers!

- Sequence Containers
  - **`std::vector`** - use for almost everything
  - **`std::deque`** - use when you need fast insertion to front AND back
- Container Adaptors
  - **`std::stack`** and **`std::queue`**
- Associative Containers
  - **`std::map`** and **`std::set`**
  - if using simple data types/you're familiar with hash functions, use **`std::unordered_map`** and **`std::unordered_set`**